

THE LANGUAGE MODELING PROBLEM



The Language Modeling Problem

- One of the oldest problems studied in statistical NLP because very useful for many applications

The Language Modeling Problem

- We have some finite vocabulary, say $V = \{\text{the, a, man, telescope, Beckham, two, ...}\} \Rightarrow$ usually several tens of thousands
- We have an (infinite) set of strings, V^+ \Rightarrow set of all possible sentences in this language
- A sentence must have 0 or more words and each word must come from V , any sequence is possible.

The Language Modeling Problem

- Possible sentences :
 - The STOP
 - A STOP
 - The fan STOP
 - The fan saw Beckham STOP
 - The fan saw saw STOP
 - The the the STOP
 - STOP

The Language Modeling Problem

- We have a *training sample* of example sentences in English
- Collection of sentences from the New York Times during the last ten years for example,
- or large sample of sentences from the web
- 90s => 20 million words
- 2000s => 1 billion words
- Nowadays => 100s billions words

The Language Modeling Problem

- Our task is to « learn » a probability distribution p over the sentences in our language.
- 2 conditions :
 - $p(x) \geq 0 \forall x \in V^+ \Rightarrow$ For any sentence x , the probability of that sentence must be greater or equal to 0
 - $\sum_{(x \in V^+)} p(x) = 1 \Rightarrow$ If we sum over all of the probabilities of the sentences in the language we obtain 1, meaning p is a well-formed distribution.

The Language Modeling Problem


- p is essentially a function which returns the probability for a sequence in a given language.
- $P(\text{the STOP}) = 10^{-12}$
- $P(\text{the fan STOP}) = 10^{-8}$
- $P(\text{the fan saw Beckham STOP}) = 2 \times 10^{-8}$
- $P(\text{the fan saw saw}) = 10^{-15}$
- ... assign a probability to every sequence in the language

The Language Modeling Problem

- We want to try and assign a high probability to likely sentences in English and low probability to unlikely sentences in English

Why would we want to do this ?!

- Language models are useful in many applications:
 - Speech recognition: language models are critical for modern speech recognizers (handwriting recognition also)
 - The estimation techniques used for this problem are useful for other NLP problems such as POS tagging or automatic translation.



Language modeling for Speech Recognition

- Quick sketch :
 - Input => an acoustic recording
 - Then map this input to the words which are actually spoken

Language modeling for Speech Recognition

- Imagine the person says « recognize speech »
- In practice, there are actually many alternative sentences which could have been spoken :
 - « wreck a nice beach »
 -
 -
- Similar sentences from an acoustic point of view

Language modeling for Speech Recognition

- A language model allows us to produce a probability for each sentence and estimate that « recognize speech » is more probable than another option.
- => Adds some very useful info to get rid of these kinds of confusions

A naive method for Language Modeling

- We have N sentences
- For any sentence or sequence $x_1 \dots x_n$, $c(x_1 \dots x_n)$ is the number of times the sentence was seen in our training data.
- A naive estimate :
- $$p(x_1 \dots x_n) = \frac{c(x_1 \dots x_n)}{N}$$

A naive method for Language Modeling

- Has some deficiencies, although it's a well-formed language model:
- Mainly it assigns proba 0 to any sentence not seen in our training sample...
- Cannot generalize to new sentences

Trigram Models

- Widely used language model
- Build heavily on the idea of Markov processes...

Markov Processes

- Consider a sequence of random variables X_1, \dots, X_n .
- Each random variable can take any value in a finite set V (vocab).
- We can assume the length n is fixed for now. ($n=100$ for ex.)
- We want to model the joint probability

$$P(X_1 = x_1, \dots, X_n = x_n)$$

Markov Processes

- Huge number of possible values.
- $|V|^n$ possible sequences in our example

First-Order Markov Process

- Going to use chain rule to decompose this joint proba
- Remember: $P(A,B) = P(A) \times P(B|A)$
- And therefore $P(A,B,C) = P(A) \times P(B|A) \times P(C|A,B)$
- So :

$$\begin{aligned} &P(X_1 = x_1, \dots, X_n = x_n) \\ &= \\ &P(X_1 = x_1) \prod_{i=2}^n P(X_i = x_i | X_1 = x_1, \dots, X_{i-1} = x_{i-1}) \end{aligned}$$

First-Order Markov Process

- The 1st order Markov assumption states that

$$\prod_{i=2}^n P(X_i = x_i | X_1 = x_1, \dots, X_{i-1} = x_{i-1})$$

Is equal to

$$\prod_{i=2}^n P(X_i = x_i | X_{i-1} = x_{i-1})$$

First-Order Markov Process

$$P(X_1 = x_1, \dots, X_n = x_n)$$

(exact equality)

$$P(X_1 = x_1) \prod_{i=2}^n P(X_i = x_i | X_1 = x_1, \dots, X_{i-1} = x_{i-1})$$

(Markov assumption)

$$P(X_1 = x_1) \prod_{i=2}^n P(X_i = x_i | X_{i-1} = x_{i-1})$$

First-Order Markov Process

- Huge assumption to state that the probability of a word here is only conditioned on the previous word...

Second-Order Markov Processes

- Very similar model :

$$P(X_1 = x_1, \dots, X_n = x_n)$$

$$= P(X_1 = x_1) P(X_2 = x_2 | X_1 = x_1) \prod_{i=3}^n P(X_i = x_i | X_{i-2} = x_{i-2}, X_{i-1} = x_{i-1})$$

- Condition on previous 2 elements vs only the previous element

Second-Order Markov Processes

- We would also like to make the length of a sentence be a random variable: not all sentences will have 100 words...
- So we can define X_n to always be equal to STOP where STOP is a special symbol.
- Basically, if STOP is at position i , then this marks the end of the sentence and $i = n$

Trigram Language Model

- Given these concepts we can define a trigram language model, which consists of :
- A finite set V
- A parameter $q(w|u, v)$ for each trigram u, v, w such that $w \in V \cup \{STOP\}$ and $u, v \in V \cup \{*\}$ (special start symbols)

Trigram Language Model

Formal Definition

- For any sentence x_1, \dots, x_n
 - where $x_i \in V$ for $i = 1, \dots, n - 1$
 - and $x_n = \text{STOP}$
- The probability of the sentence under the trigram language model is

$$p(x_1, x_2, \dots, x_n) = \prod_{i=1}^n q(x_i | x_{i-2}, x_{i-1})$$

- Where we define $x_{-1} = x_0 = *$

An example to make things clearer

- Sentence : « * * The dog barks STOP »
 - $p(** the dog barks STOP) = q(the|*,*) \times q(dog|*,the) \times q(barks|the,dog) \times q(STOP|dog,barks)$
- Product of terms to get the proba of the sentence under this type of language model
- We're treating sentences as being generated by a second order Markov process, where each word generated is dependent purely on the 2 previous words.

Trigram Language Model

- Advantages :
 - Simple, easy and cheap
 - useful for many applications
 - availability of statistics over the internet
 - well understood math
- Disadvantages:
 - Language: they do not capture non-local dependencies

Estimating the parameters

- So we need to estimate $q(w_i|w_{i-2}, w_{i-1})$
- Remember, if we have two dependent events :

$$p(A, B) = p(A) \times p(B|A)$$

- Which is equivalent to

$$p(B | A) = \frac{p(A, B)}{p(A)}$$

- Which can be generalized to 3 events

$$p(C|A, B) = \frac{p(A, B, C)}{p(A, B)}$$

Estimating the parameters

- A natural estimate is therefore :

$$q(w_i | w_{i-2}, w_{i-1}) = \frac{\text{Count}(w_{i-2}, w_{i-1}, w_i)}{\text{Count}(w_{i-2}, w_{i-1})}$$

- So for example :

$$q(\text{laughs} | \text{the}, \text{dog}) = \frac{\text{Count}(\text{the}, \text{dog}, \text{laughs})}{\text{Count}(\text{the}, \text{dog})}$$

Estimating the parameters from a toy corpus

- An example corpus:

1. the cat saw the mouse.
2. the cat heard a mouse.
3. the mouse heard.
4. a mouse saw.
5. a cat saw.
6. a cat heard the mouse.

=> Using the corpus, give the parameter estimates for :

- a bigram language model
- a trigram language model

(the parameters for the first 2 sentences are enough)

Estimating the parameters

Bigram	Count	Unigram	Count	Relative frequency
* the	3	*	6	3/6
the cat	2	the	5	2/5
cat saw	2	cat	4	2/4
saw the	1	saw	3	1/3
the mouse	2	the	5	2/5
mouse STOP	3	mouse	5	3/5
cat heard	2	cat	4	2/4
heard a	1	heard	3	1/3
a mouse	2	a	4	2/4
...

Estimating the parameters

Trigrams	Count	Bigram	Count	Relative frequency
* * the	3	**	6	3/6
* the cat	2	* the	3	2/3
the cat saw	1	the cat	2	1/2
cat saw the	1	cat saw	2	1/2
saw the mouse	1	saw the	1	1
the mouse STOP	2	the mouse	3	2/3
the cat heard	1	the cat	2	1/2
cat heard a	1	cat heard	2	1/2
heard a mouse	1	heard a	1	1
a mouse STOP	1	a mouse	2	1/2
...