

Linear Regression.

A number of slides and screenshots from : [Andrew Ng's](#) course on machine learning and [Sebastian Raschka's](#) course on deep learning
Both can be found for free on youtube

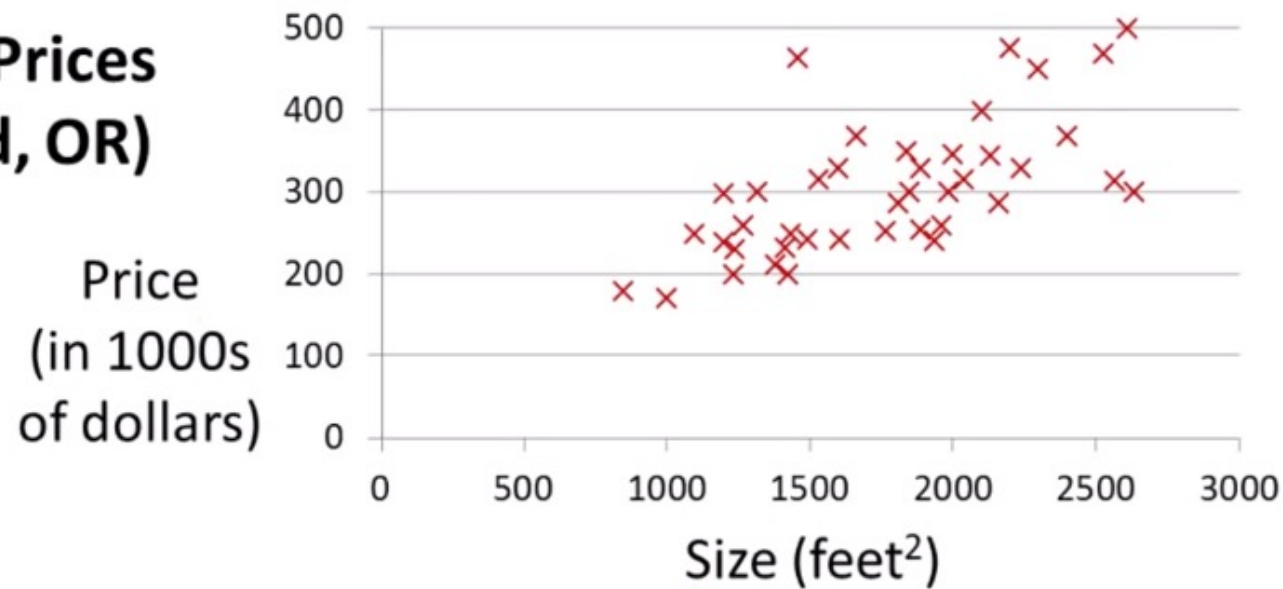
What does regression mean ?

- Seen in intro, but :
- Regression means predicting real-valued outputs.
- An essential type of supervised machine learning task : for each example in the data, we want to get as close as possible to the real-valued label.
- Often contrasted with classification (discrete labels).
- Example :
 - Predicting height => many many real-valued outputs are possible...
 - Vs. Predicting a « height class » : short | medium-height | tall

Dataset and problem example

- Imagine we want to create an ML algorithm to predict the price of a house, using only as information the size of the house. This is the dataset we can use to train our algorithm.

Housing Prices (Portland, OR)



Training Set and Notation

| Training set of housing prices (Portland, OR) | Size in feet ² (x) | Price (\$) in 1000's (y) |
|---|-----------------------------------|------------------------------|
| | 2104 | 460 |
| | 1416 | 232 |
| | 1534 | 315 |
| | 852 | 178 |
| | ... | ... |

Notation:

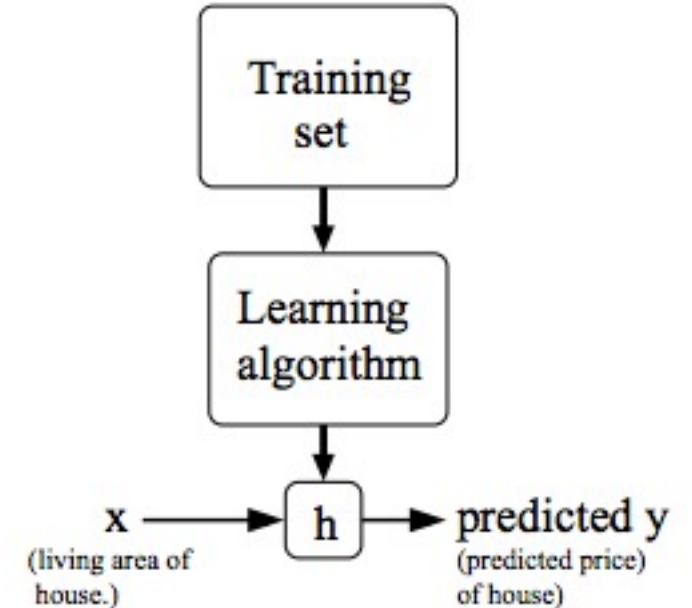
m = Number of training examples

x 's = "input" variable / features

y 's = "output" variable / "target" variable

The supervised learning workflow

- h : hypothesis
- h is a function which maps x 's to y 's
- Our goal will be to find the function which takes x as input and predicts the correct y for that x .

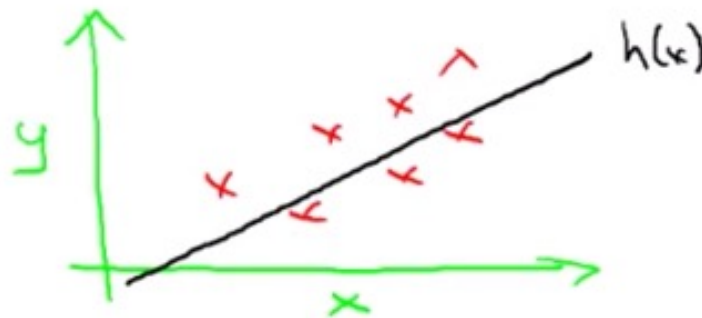


Model function h

- To start with, we will use a simple model, a function which is the equation of a line (maybe you remember $y = ax + b$ from school ?)

$$h(x) = \theta_0 + \theta_1 x$$

- This model will predict that y is some linear (straight line) function :



If this seems a bit odd to you...

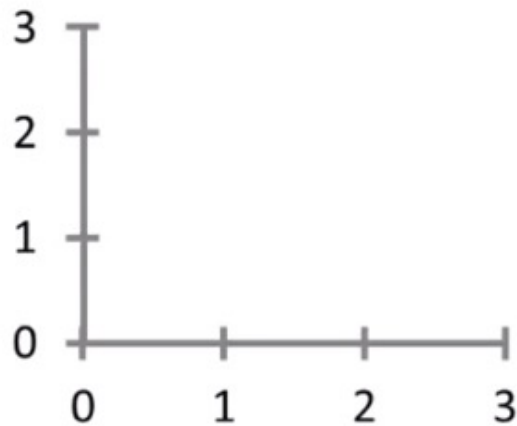
- Remember we want our function to predict the examples we have in our training set correctly,
- which our simple model will probably not do very well....
- What if we can't get to all the points using a straight line ?
- Don't worry for now, this is still a good starting point !

Cost function

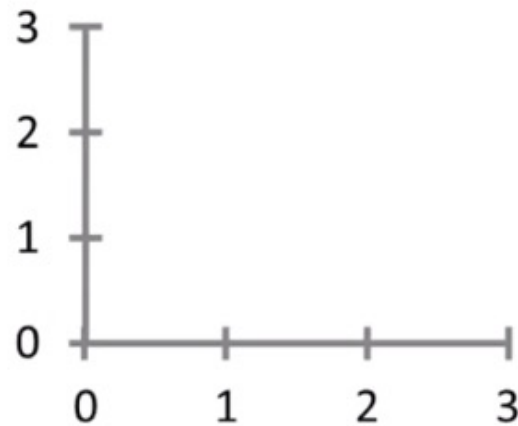
- This is a **second** function we will use to judge how well our straight line fits the data and to find the best possible straight line.
- $h(x) = \theta_0 + \theta_1 x \Rightarrow$ our model
- θ_i are what we call **parameters** and we want to find the right combination of those parameters to get the best line.
- So **how do we choose the right parameters ?**

Different parameter choices/hypotheses

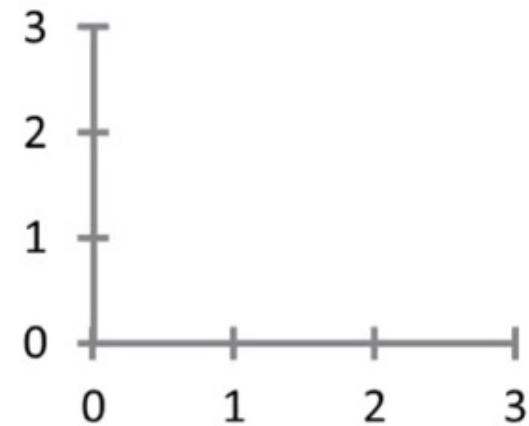
$$h_{\theta}(x) = \theta_0 + \theta_1 x$$



$$\theta_0 = 1.5$$
$$\theta_1 = 0$$



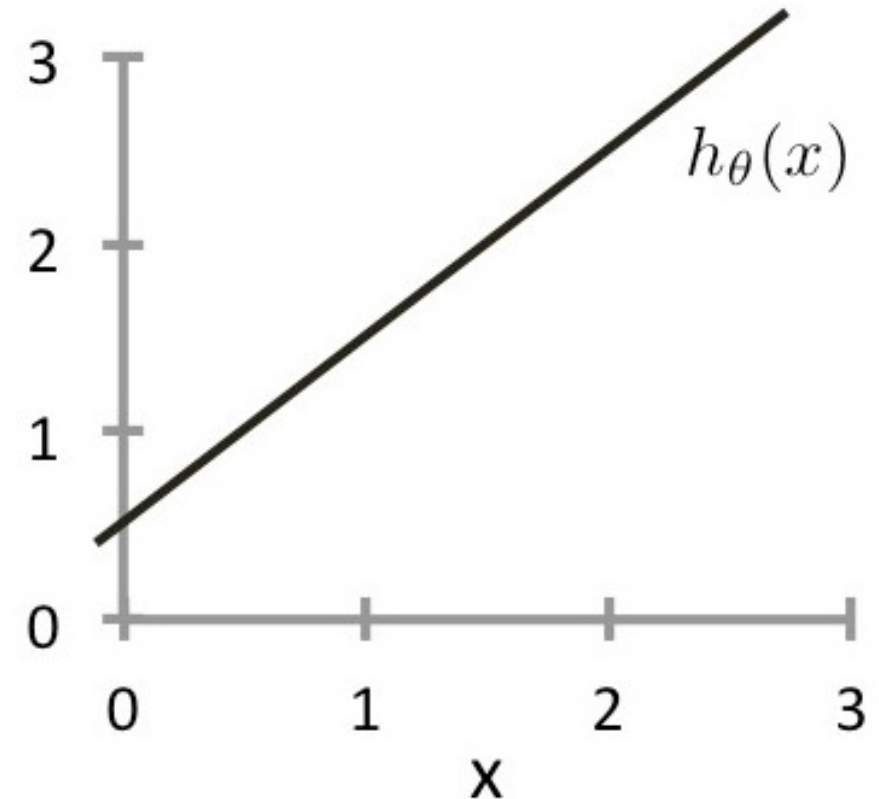
$$\theta_0 = 0$$
$$\theta_1 = 0.5$$



$$\theta_0 = 1$$
$$\theta_1 = 0.5$$

Exercise

- Look at the plot of $h(x) = \theta_0 + \theta_1 x$
- What are the values of θ_0 and θ_1 ?



Minimization Problem

- We want to choose θ_0 and θ_1 so that
- $h(x)$ is close to y for our training examples (x, y) ...
- So this is actually a **minimization problem**,
- where we want to minimize $(h(x) - y)^2$ for example, by tweaking our parameters θ_0 and θ_1

Cost function = Quantifying the model's error

- For all of our examples m the average error is :

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$

The 2 is just there to make the math easier but doesn't change anything fundamentally, you can regard this as the average error.

- This function is known as the MSE (we'll see how it works in a few slides) and is the most commonly used:

Mean Squared Error

To recap

Hypothesis:

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Parameters:

$$\theta_0, \theta_1$$

Cost Function:

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Goal: minimize $J(\theta_0, \theta_1)$
 θ_0, θ_1

Cost Function Intuition

- Let's use a simplified model hypothesis to understand what's going on:

$$h(x) = \theta_1 x$$

- Our objective is now to minimize

$$J(\theta_1)$$

- And our cost function looks like

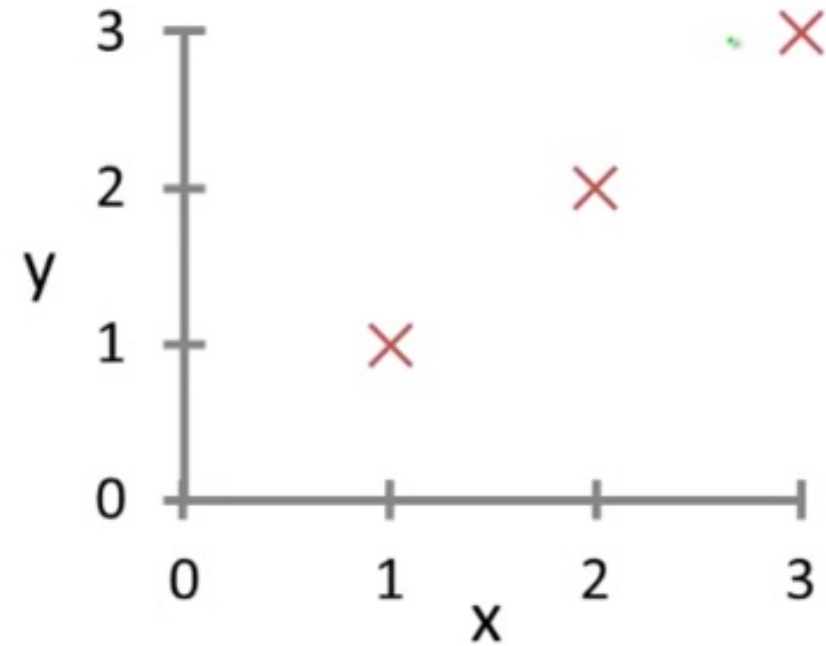
$$\frac{1}{2m} \sum_{i=1}^m (\theta_1 x^i - y^i)^2$$

Hypothesis function vs. Cost function

- If the points on the graph represent our training data and $\theta_1 = 1$, what does our hypothesis (line) look like ?

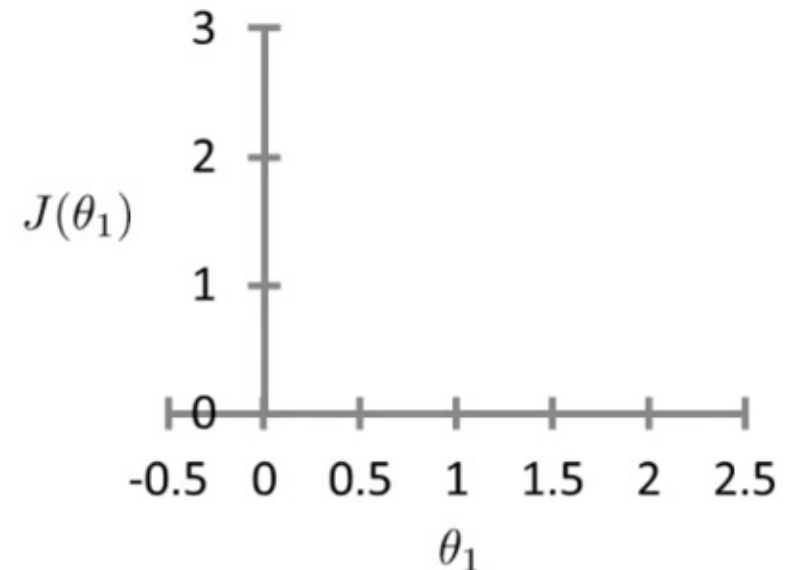
- What is the cost ?

- Remember : $\frac{1}{2m} \sum_{i=1}^m (\theta_1 x^i - y^i)^2$



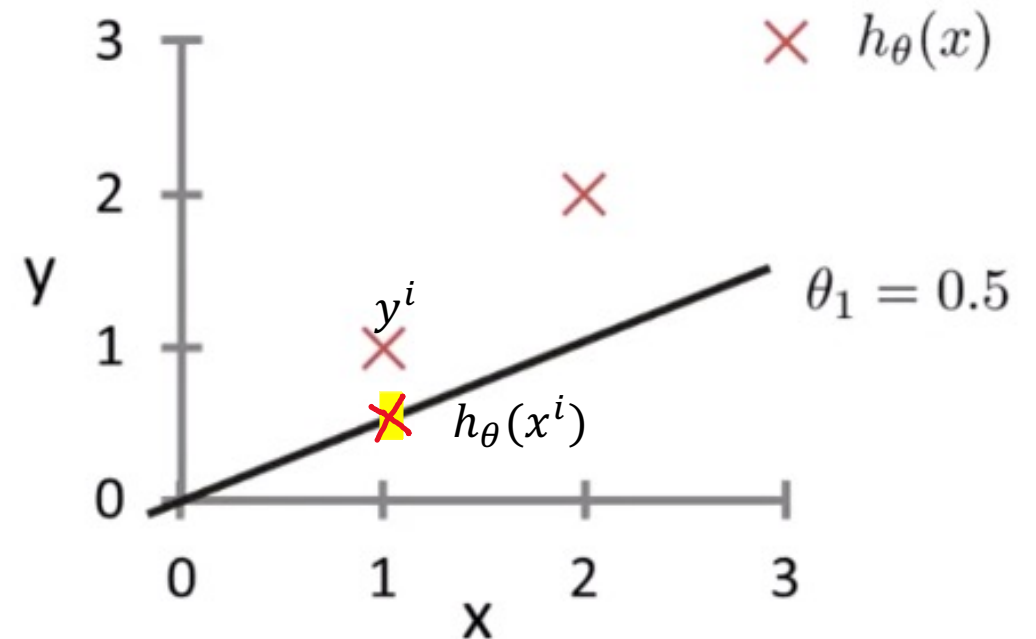
Hypothesis function vs. Cost function

- $J(\theta_1 = 1) = 0$
- We can now plot our error rate
- Notice that the values for θ_1 are on the horizontal axis. This is not the same graph as before !!
- This is a plot for the **cost function** :



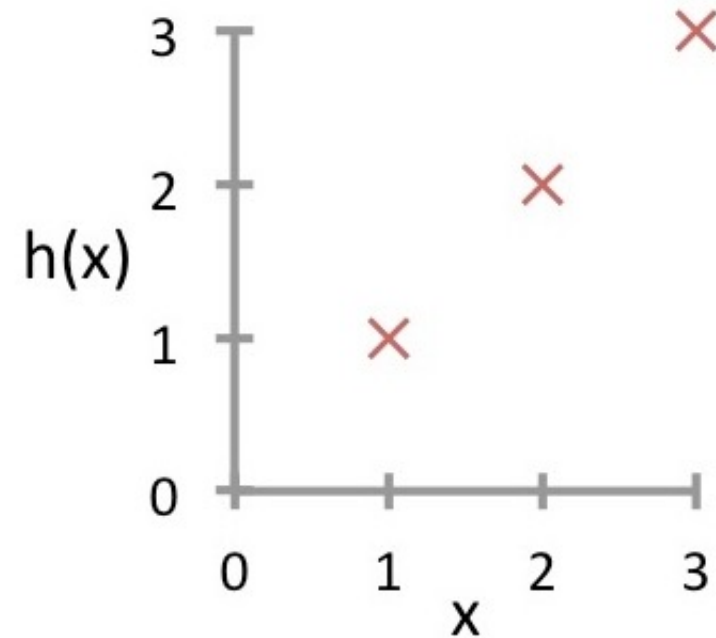
Hypothesis function vs. Cost function

- Now let's look at $\theta_1 = 0.5$
- And compute $J(\theta_1 = 0.5)$ (approx. 0.58)
- The error for each point is actually the height which separates the data point and the line for a given x .



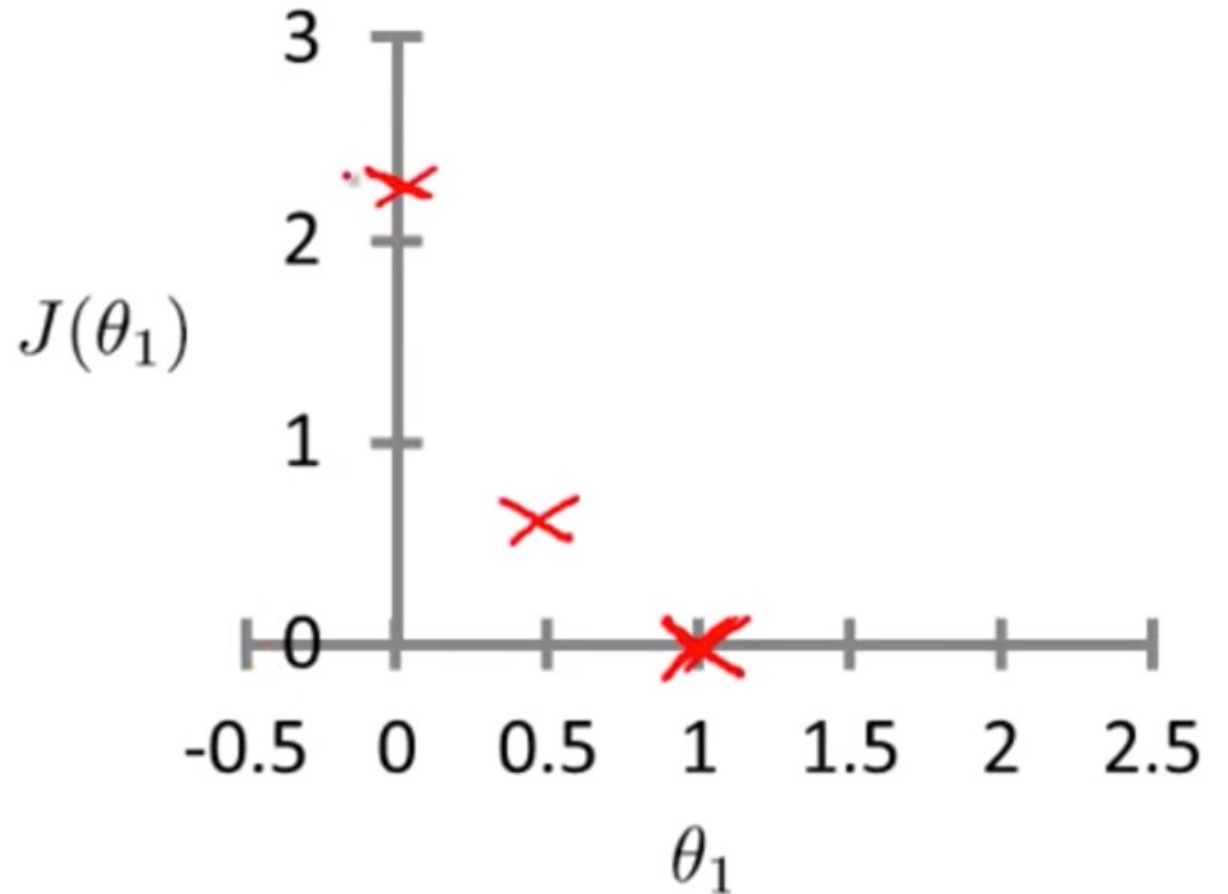
Your turn !

- Suppose this is our training set. $m = 3$.
- Given the same hypothesis and cost functions as before, what is $J(0)$?
- ie. $\theta_1 = 0$
- Should be approx. 2.3



Hypothesis function vs. Cost function

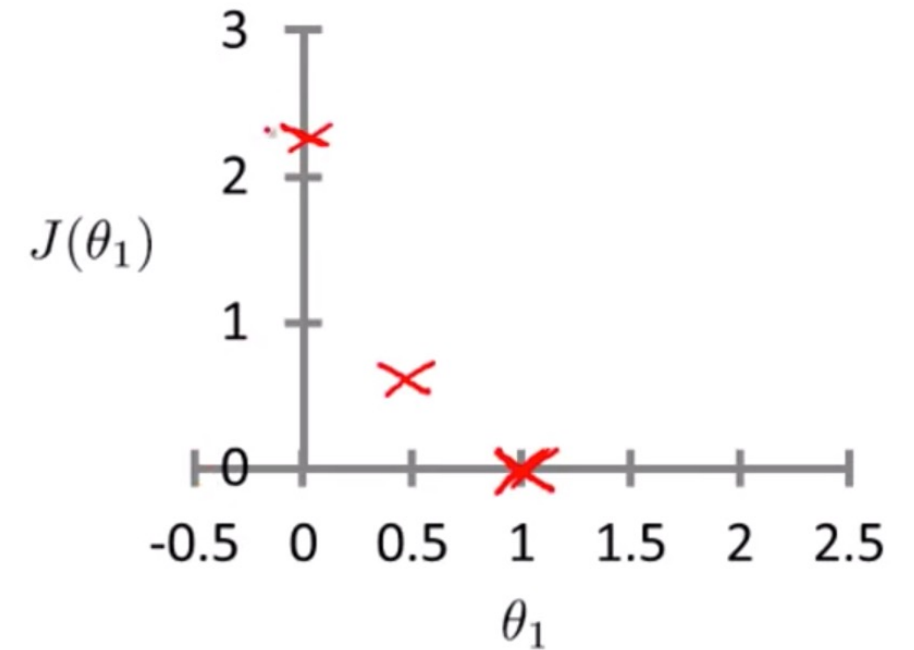
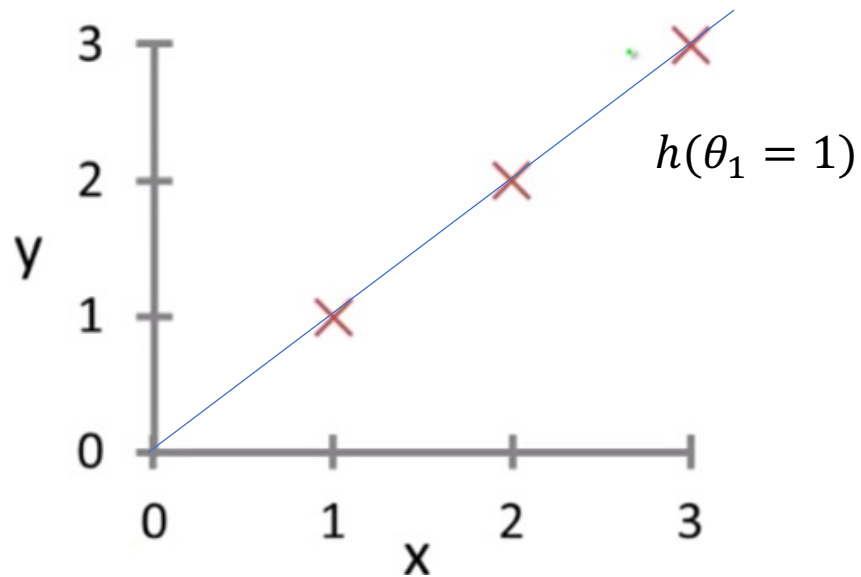
- We could continue plotting points but we'll stop here.
- With the error calculated for the different values of θ_1 , we start to see part of the general shape of the function
- It turns out the function is convex/looks like a parabola.



Quick recap

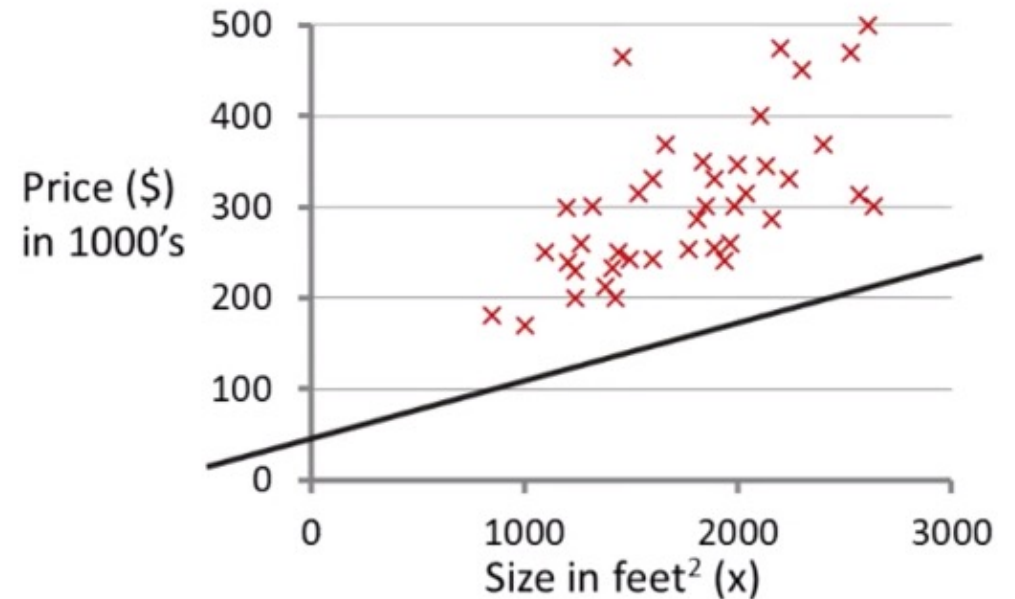
- Each value of θ_1 plotted corresponds to a different hypothesis / model on the data point graphs shown previously.
- For each value of θ_1 we can compute a value $J(\theta_1)$ to trace out the cost function.
- Now remember, we wanted to find the value of θ_1 which minimized $J(\theta_1)$... Looking at the graph we can now do so !

- No surprise, the value of θ_1 which minimizes the error, is associated with the model which fits the data perfectly



Back to 2 parameters

- Now, going back to our original data and model, we use a 2 parameter hypothesis to draw our line.
- For :
- $\theta_0 = 50$
- $\theta_1 = 0.06$
- We get this straight line as our model



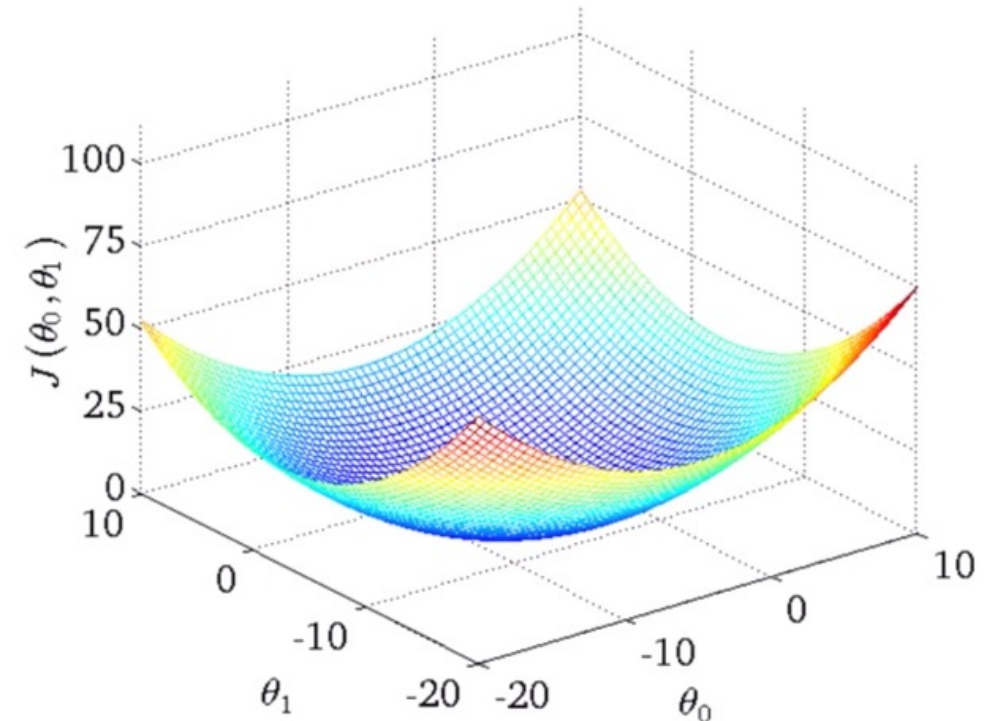
$$h_{\theta}(x) = 50 + 0.06x$$

Corresponding Cost function

- Now we have two parameters, the error graph will be slightly harder to plot as it has 3 dimensions:

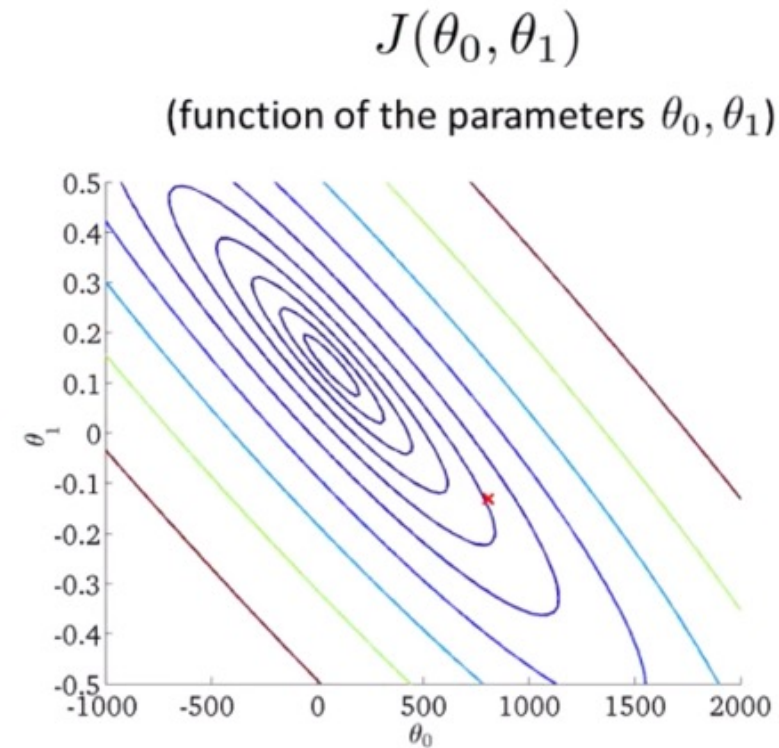
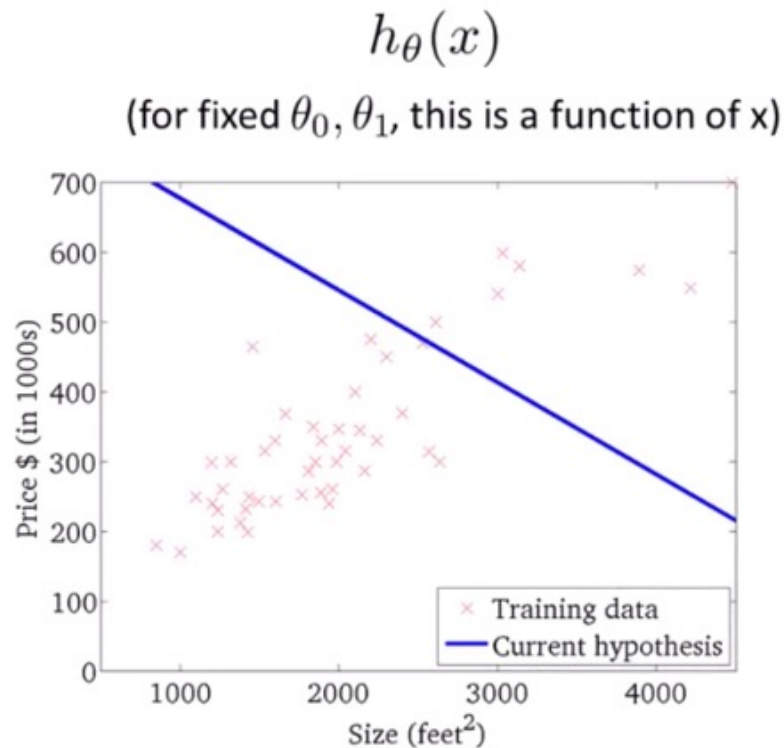
$$\theta_1, \theta_2, cost$$

- Indeed , $J(\theta_1, \theta_2)$ now has 2 inputs,
- So it will look like this in 3D:



Contour Plots

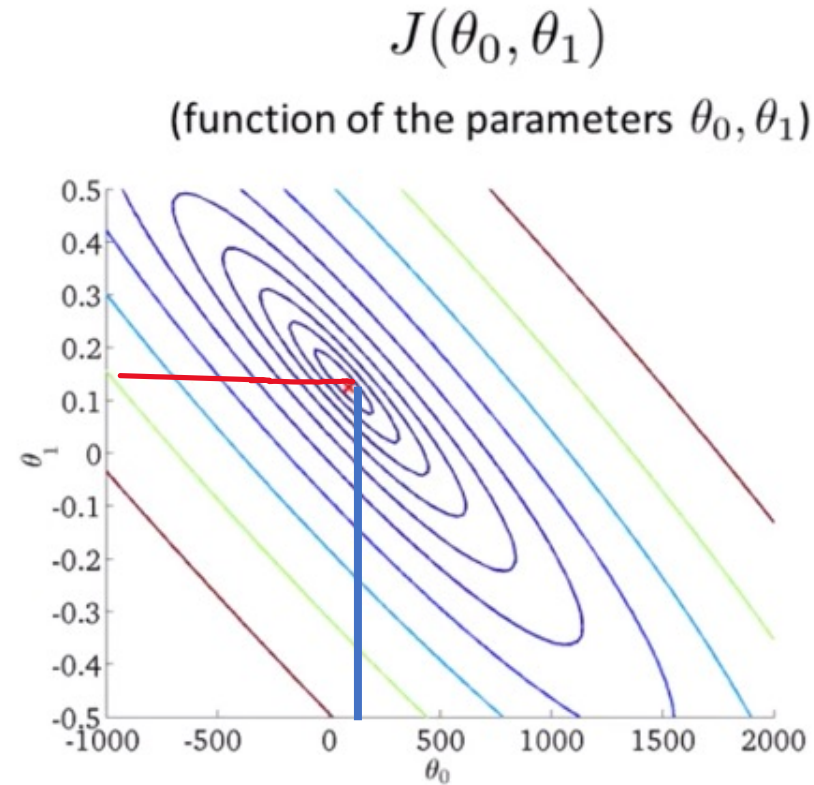
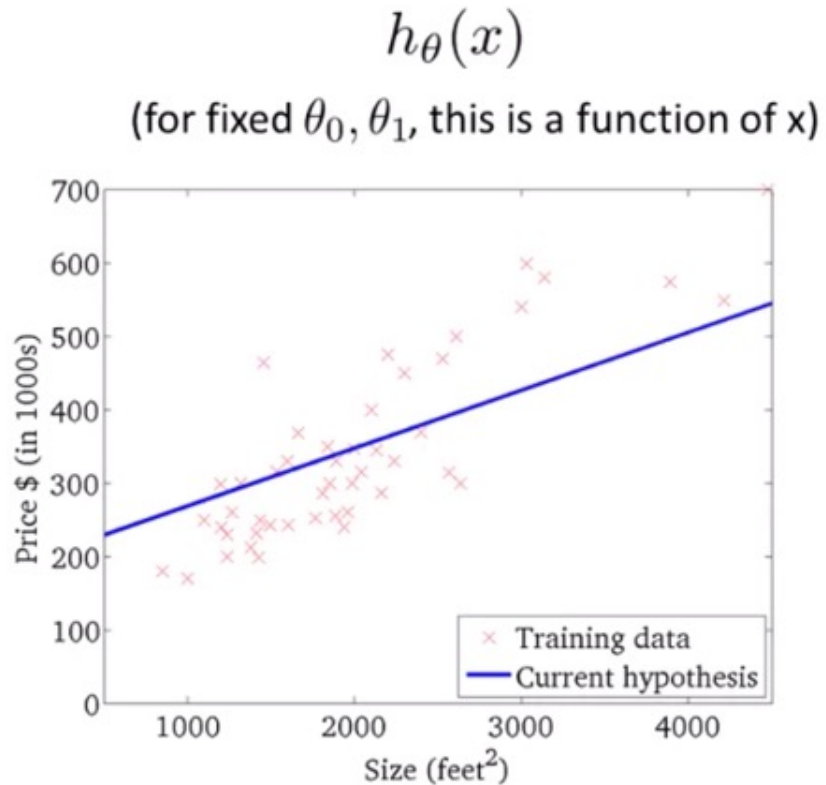
- To stay in 2D, you will see the cost function represented by a contour plot :



The ovals/ellipses show the set of points which take on the same value for given values of θ_0, θ_1

Countour Plots

- The minimum is at the center of all the « ellipses ».
- This shows a model very close to the minimum.



Gradient Descent

- Now we know how to evaluate a model, using a cost function, how do we make the model *learn* the optimal parameters ?
- In other words, how do we minimize the cost function without testing all the different possible models ?
- The algorithm used to do this is called *Gradient Descent*, and is essential to most machine learning algorithms, not just linear regression !
- In DL libraries this type of algorithm is called an 'optimizer' and other variants exist.

Gradient Descent

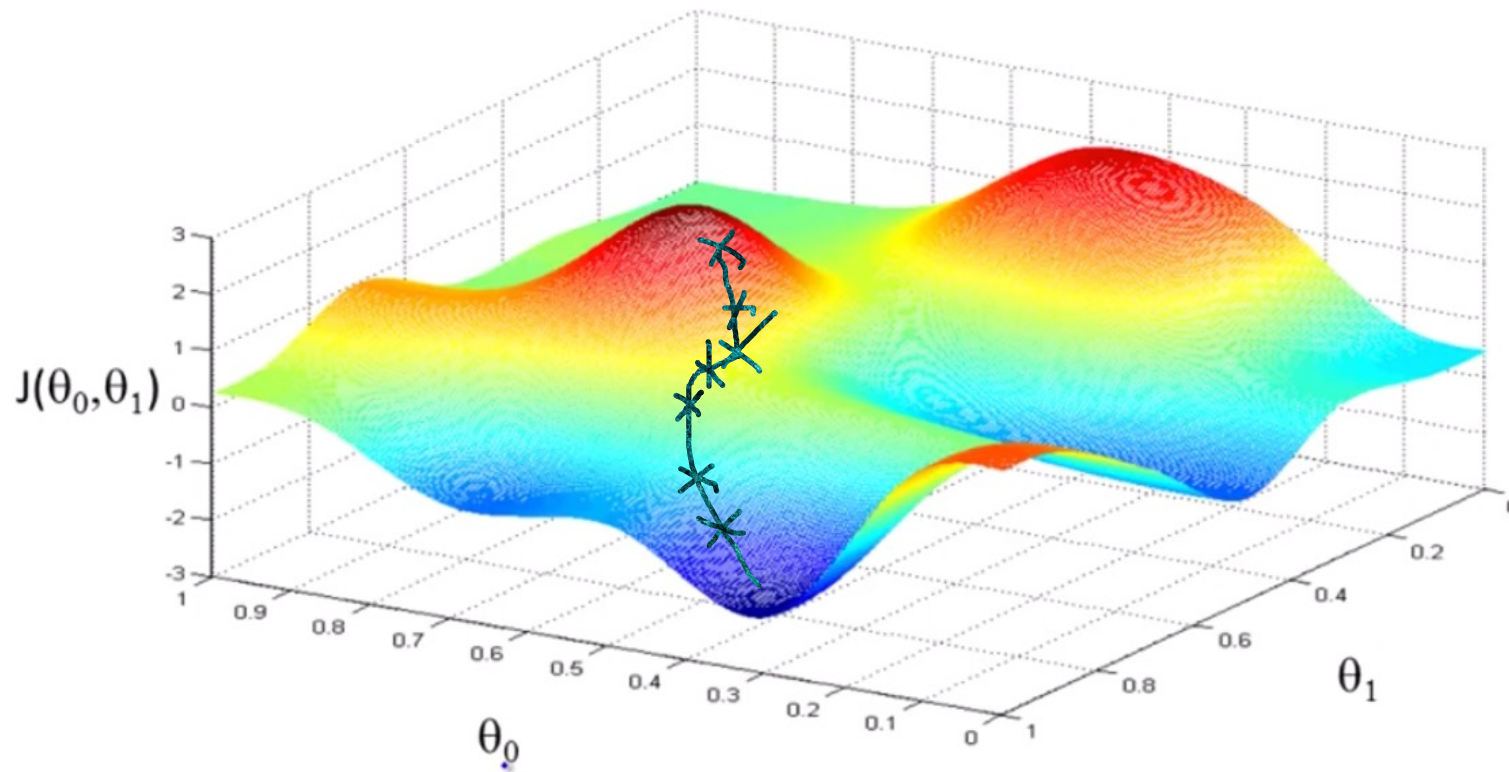
- We have some function $J(\theta_1, \theta_2)$ which we want to minimize...

Outline :

- Start with some initial guess, some random values for θ_1, θ_2
- Keep updating θ_1, θ_2 a little bit to reduce $J(\theta_1, \theta_2)$ until we hopefully end up at a minimum

GD intuition

- This is your cost function in 3D
- Imagine you start somewhere near the top of one of the « hills » and your goal is to walk in the direction which will take you down to the bottom the fastest.



GD formula

$$\begin{array}{l} \text{repeat until convergence } \{ \\ \quad \theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \quad (\text{for } j = 0 \text{ and } j = 1) \\ \} \end{array}$$

- This is the update formula for each of the parameters
- $:=$ signifies assignment
- α is a number called the *learning rate*. If α is very large, then it corresponds to an aggressive learning procedure and big steps being taken « downhill » and vice versa.
- $\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$ is a derivative term, for which we need to do a bit of calculus !

GD Intuition

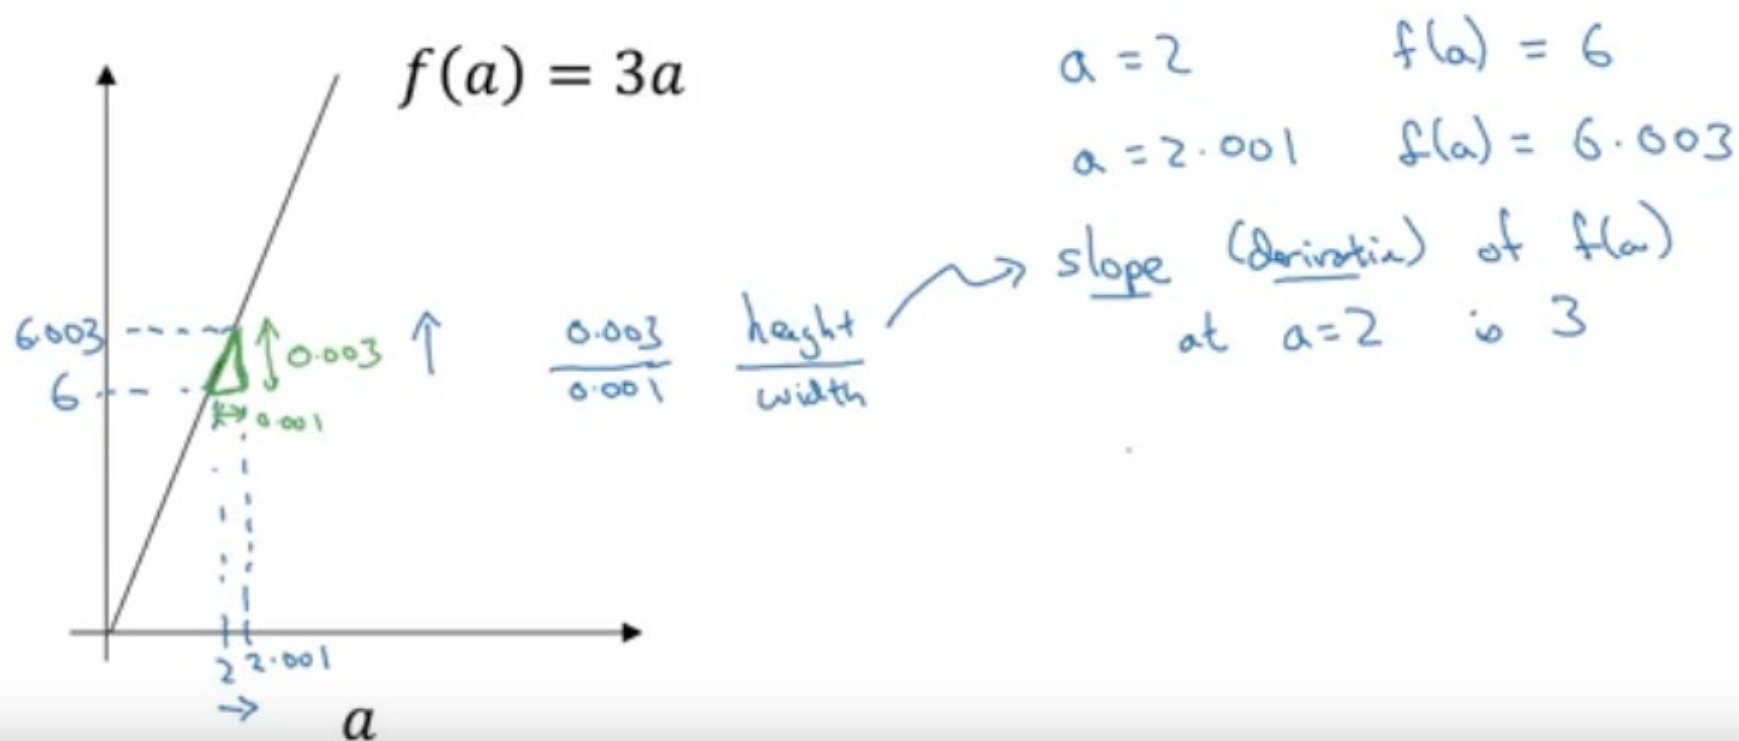
- Why does this update make sense ?
- Why are we putting those 2 terms together ?
- Let's try and get a basic understanding of derivatives before we go any further.

Derivatives

- The derivative describes how the output of a function varies with regard to a very very very tiny positive nudge to the input, to the point where we consider *almost* no variation in input....
- Informally, the derivative tells you how a function behaves at a particular « instant », i.e. for a given input value.
- The derivative is commonly referred to as an « instantaneous rate of change »

Derivatives

- Here is a linear function as an example. What happens when we shift the input by a 'small' value like 0.001
- when $a=2$?
- when $a=5$?



Derivatives

- With this function, we expect a small positive nudge in the input to make the output increase by 3 times the value of that nudge.

$$f(5.001) = 15.003$$

- In other words the **ratio** between the change in output and the change in input is 3 :

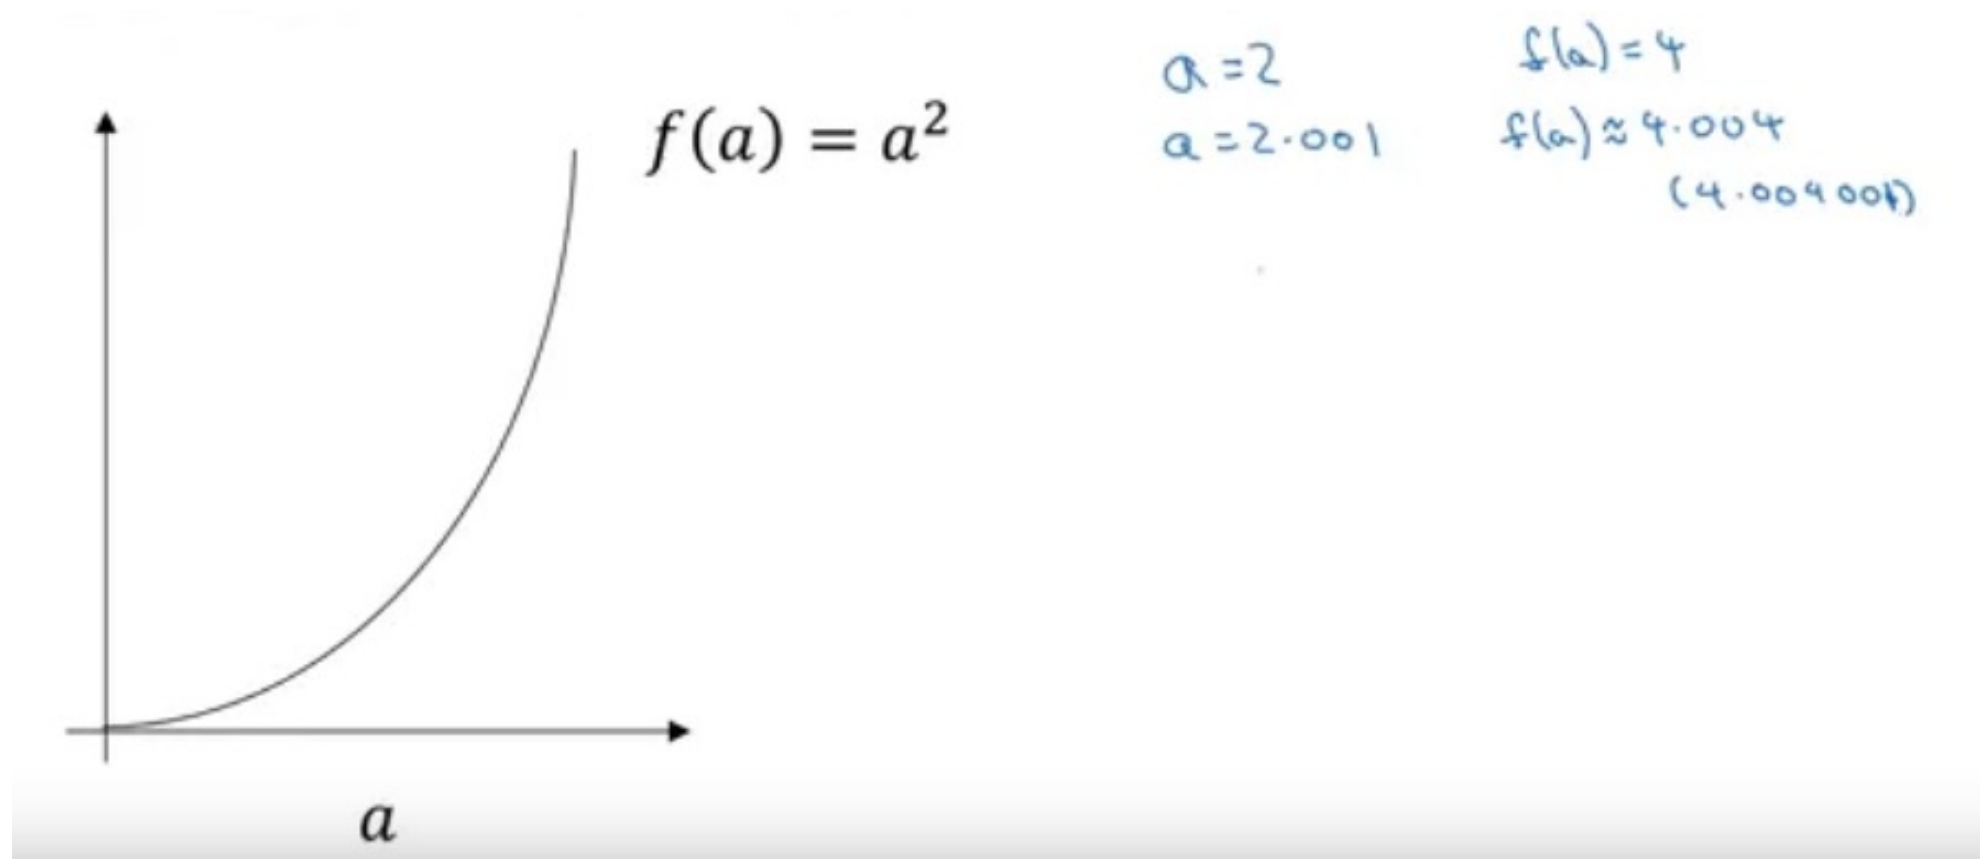
$$\frac{\text{change in } f(a)}{\text{change in } a} = \frac{df(a)}{da} = \frac{0.003}{0.001} = 3$$

Derivatives

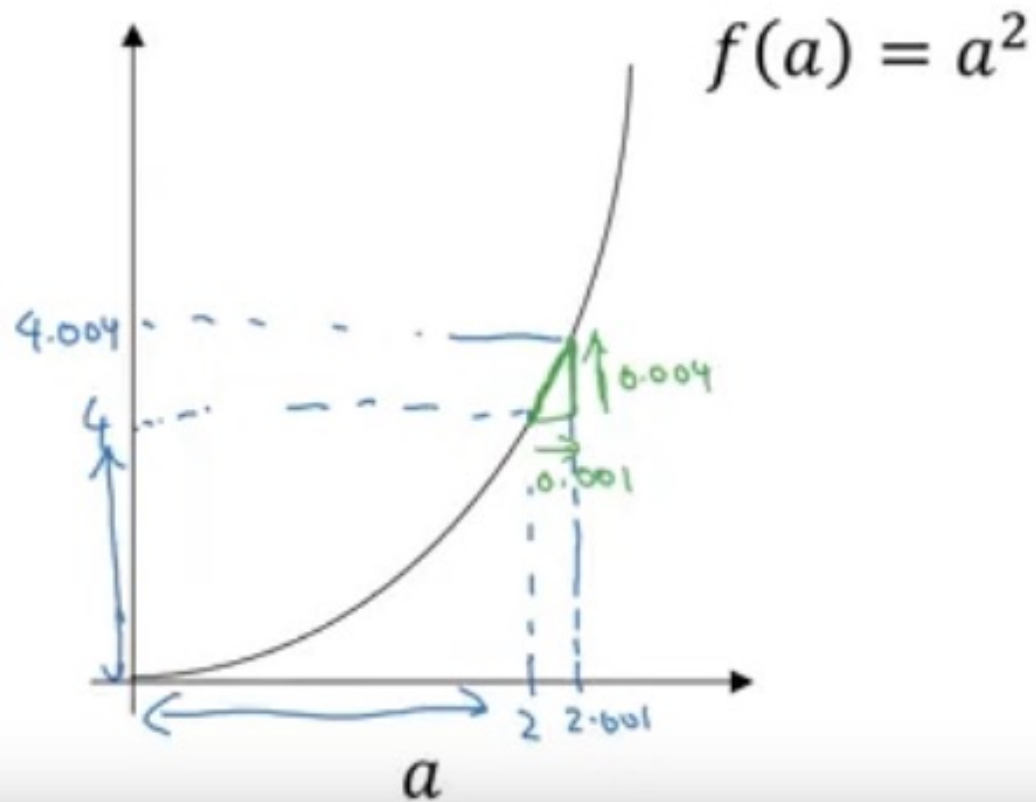
- This is just an example, but formally, the derivative considers this ratio when the input is increased by a much tinier amount !
- $\text{Nudge} < 0.000000000000.....1 \Rightarrow$ nudge gets as close to 0 as possible
- In this previous example, whatever input value we pick, the derivative will be the same.
- The little triangle will be the same for any value of a : this makes sense since the function is a line and the output increases at a constant rate
- Question : What if the derivative was negative everywhere ? What would the function look like ?

Derivatives

- What if our function isn't a line ?

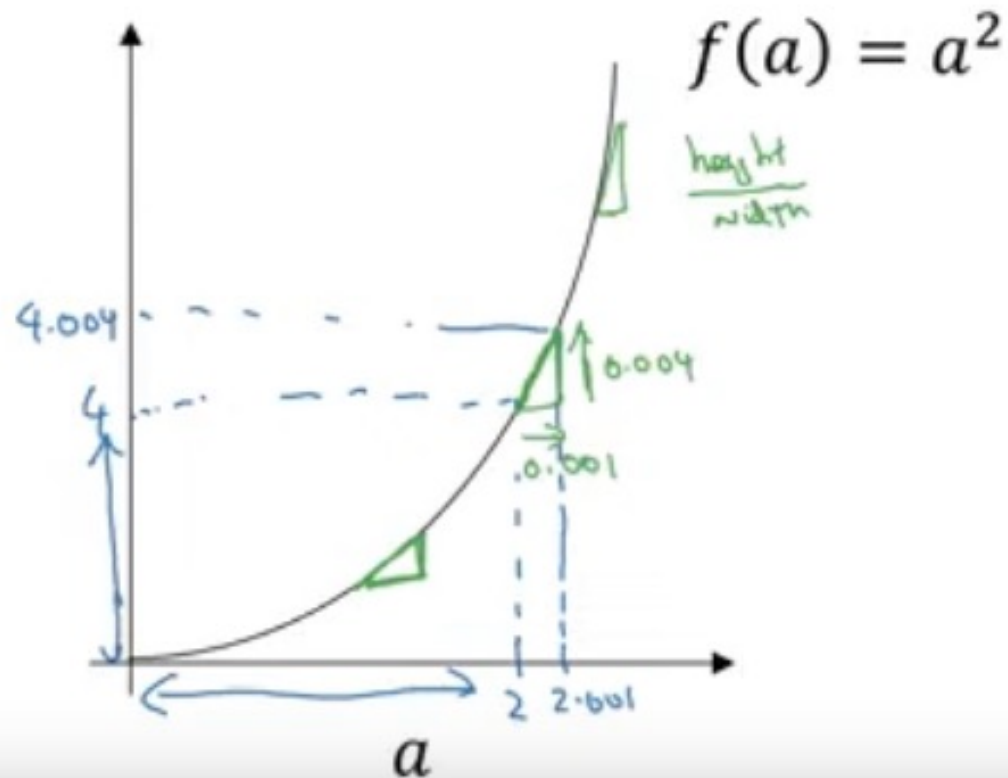


- The derivative at $a=2$ is ...



$$\begin{aligned}
 a &= 2 & f(a) &= 4 \\
 a &= 2.001 & f(a) &\approx 4.004 \\
 & & & (4.004001) \\
 \text{slope (derivative) of } f(a) \text{ at} \\
 a &= 2 & \text{ is } & 4. \\
 \frac{d}{da} f(a) &= 4 & \text{ when } a &= 2..
 \end{aligned}$$

- The derivative at $a=5$ is ...



$$\begin{aligned}
 a &= 2 & f(a) &= 4 \\
 a &= 2.001 & f(a) &\approx 4.004 \\
 & & & (4.004004) \\
 \text{slope (derivative) of } f(a) \text{ at } & & & \\
 a=2 & \text{ is } 4. \\
 \frac{d}{da} f(a) &= \underline{4} \text{ when } a = \underline{2}. \\
 a &= 5 & f(a) &= 25 \\
 a &= 5.001 & f(a) &\approx 25.010 \\
 \frac{d}{da} f(a) &= \underline{10} \text{ when } a = \underline{5}.
 \end{aligned}$$

- Rules actually exist to compute derivatives by hand quickly !
- For the function

$$f(a) = a^2$$

$$f'(a) = \frac{d}{da} f(a) = 2a$$

(The notations are called Lagrange and Leibniz notations and are both common)

- If we look at the derivatives/slopes/ratios we calculated previously, this does indeed seem to work ! (and this can be proven of course)
- Note: the derivative is equal to the slope of the tangent on the graph at our input value.

Derivatives: (optional)

$$f'(x) = \frac{df}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x + \Delta x) - f(x)}{\Delta x}$$

Example 2: $f(x) = x^2$

$$\begin{aligned} \frac{df}{dx} &= \lim_{\Delta x \rightarrow 0} \frac{f(x + \Delta x) - f(x)}{\Delta x} \\ &= \lim_{\Delta x \rightarrow 0} \frac{x^2 + 2x\Delta x + (\Delta x)^2 - x^2}{\Delta x} \\ &= \lim_{\Delta x \rightarrow 0} \frac{2x\Delta x + (\Delta x)^2}{\Delta x} \\ &= \lim_{\Delta x \rightarrow 0} 2x + \Delta x. \end{aligned}$$

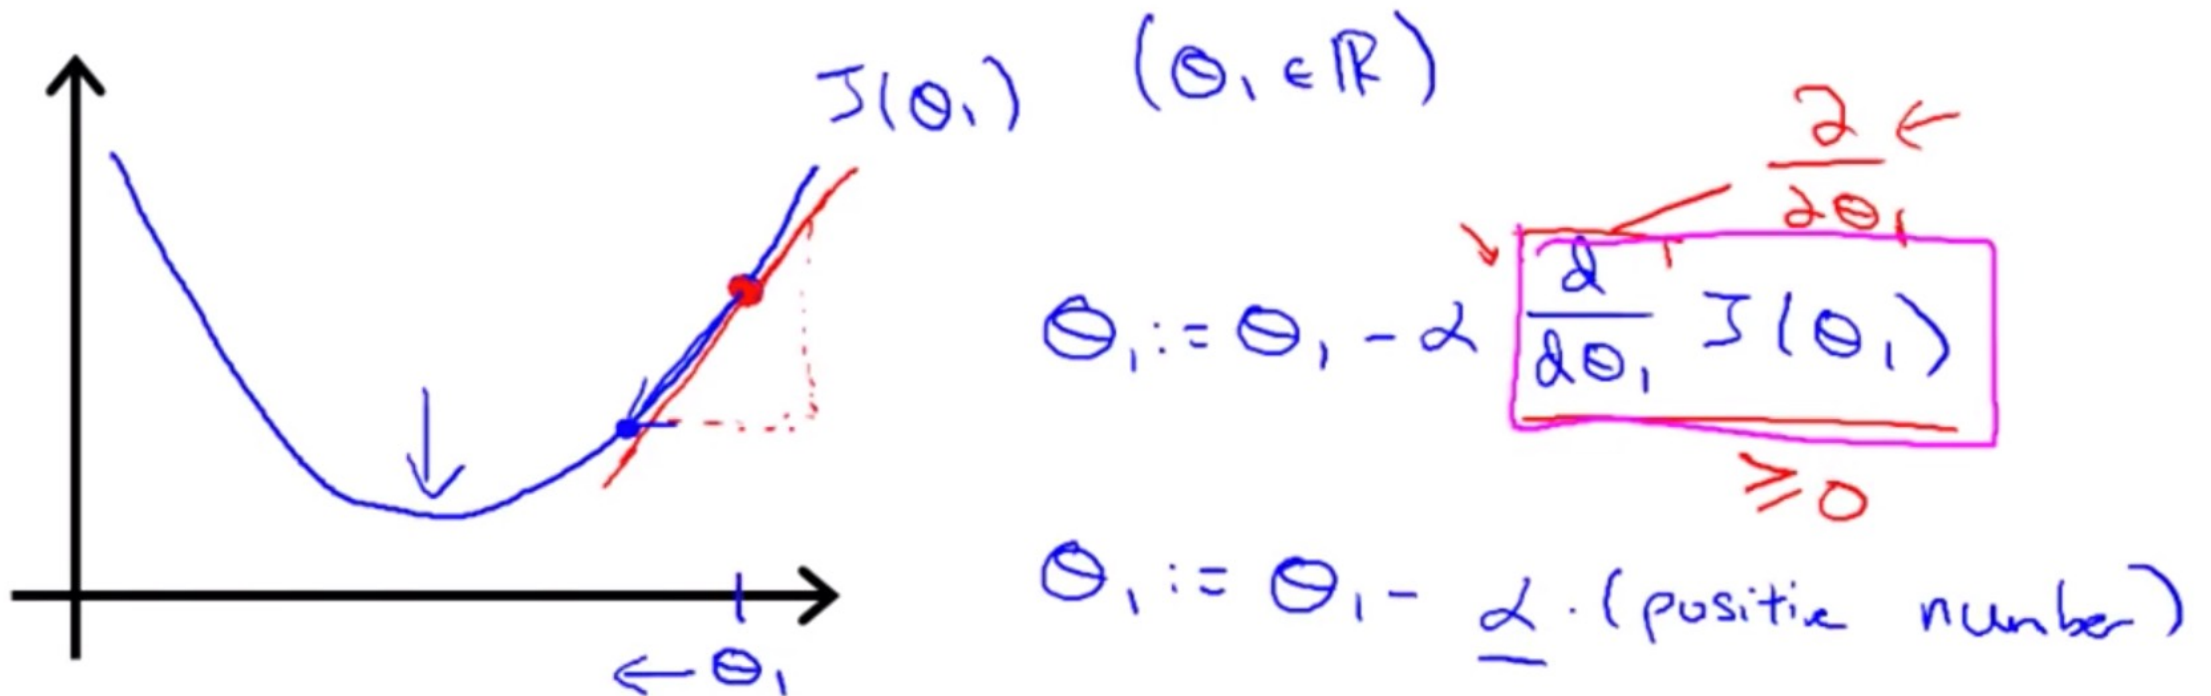
- As Δx approaches 0, the derivative
- Approaches $2x$.

GD Intuition

- Now we have a basic understanding of derivatives, let's apply this understanding to the gradient descent algorithm by using a simpler example, with a cost function of only 1 single parameter.
- We use $J(\theta_1)$ instead of $J(\theta_0, \theta_1)$
- Let's look at a couple scenarios to see how Gradient Descent updates our parameter θ_1 .

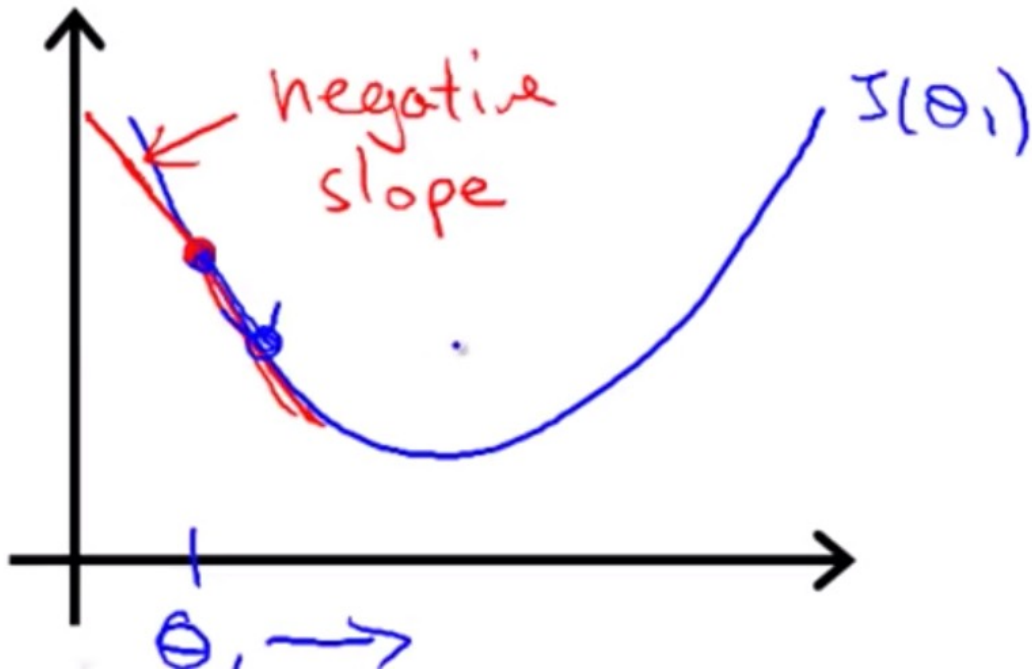
When the derivative is positive...

- Remember, our cost function looks like a parabola.
- When θ_1 is too high, we would want Gradient Descent to **reduce** this parameter and bring it closer to the « sweet spot », where the cost is minimized.
- Let's see if it does the right thing :



When the derivative is negative...

- When θ_1 is too low, let's see if Gradient Descent **increases** it and brings it closer to the « sweet spot », where the cost is minimized :



$$\frac{\partial}{\partial \theta_1} J(\theta_1) \leq 0$$
$$\theta_1 := \theta_1 - \alpha \text{ (negative number)}$$

Recap

- When the the parameter value is too high, the derivative is positive and the update rule decreases the value for the parameter.

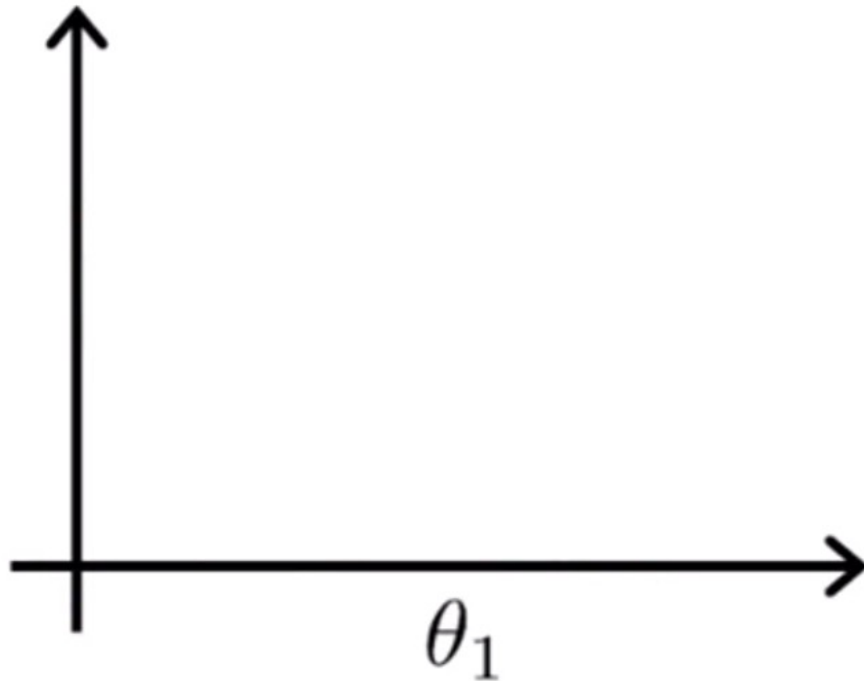
$$\theta_1 := \theta_1 - \alpha \underbrace{\frac{d}{d\theta_1} J(\theta_1)}_{> 0}$$

- Conversely, when the parameter value is too low, the parameter value will be increased by the update rule.

$$\theta_1 := \theta_1 - \alpha \underbrace{\frac{d}{d\theta_1} J(\theta_1)}_{< 0}$$

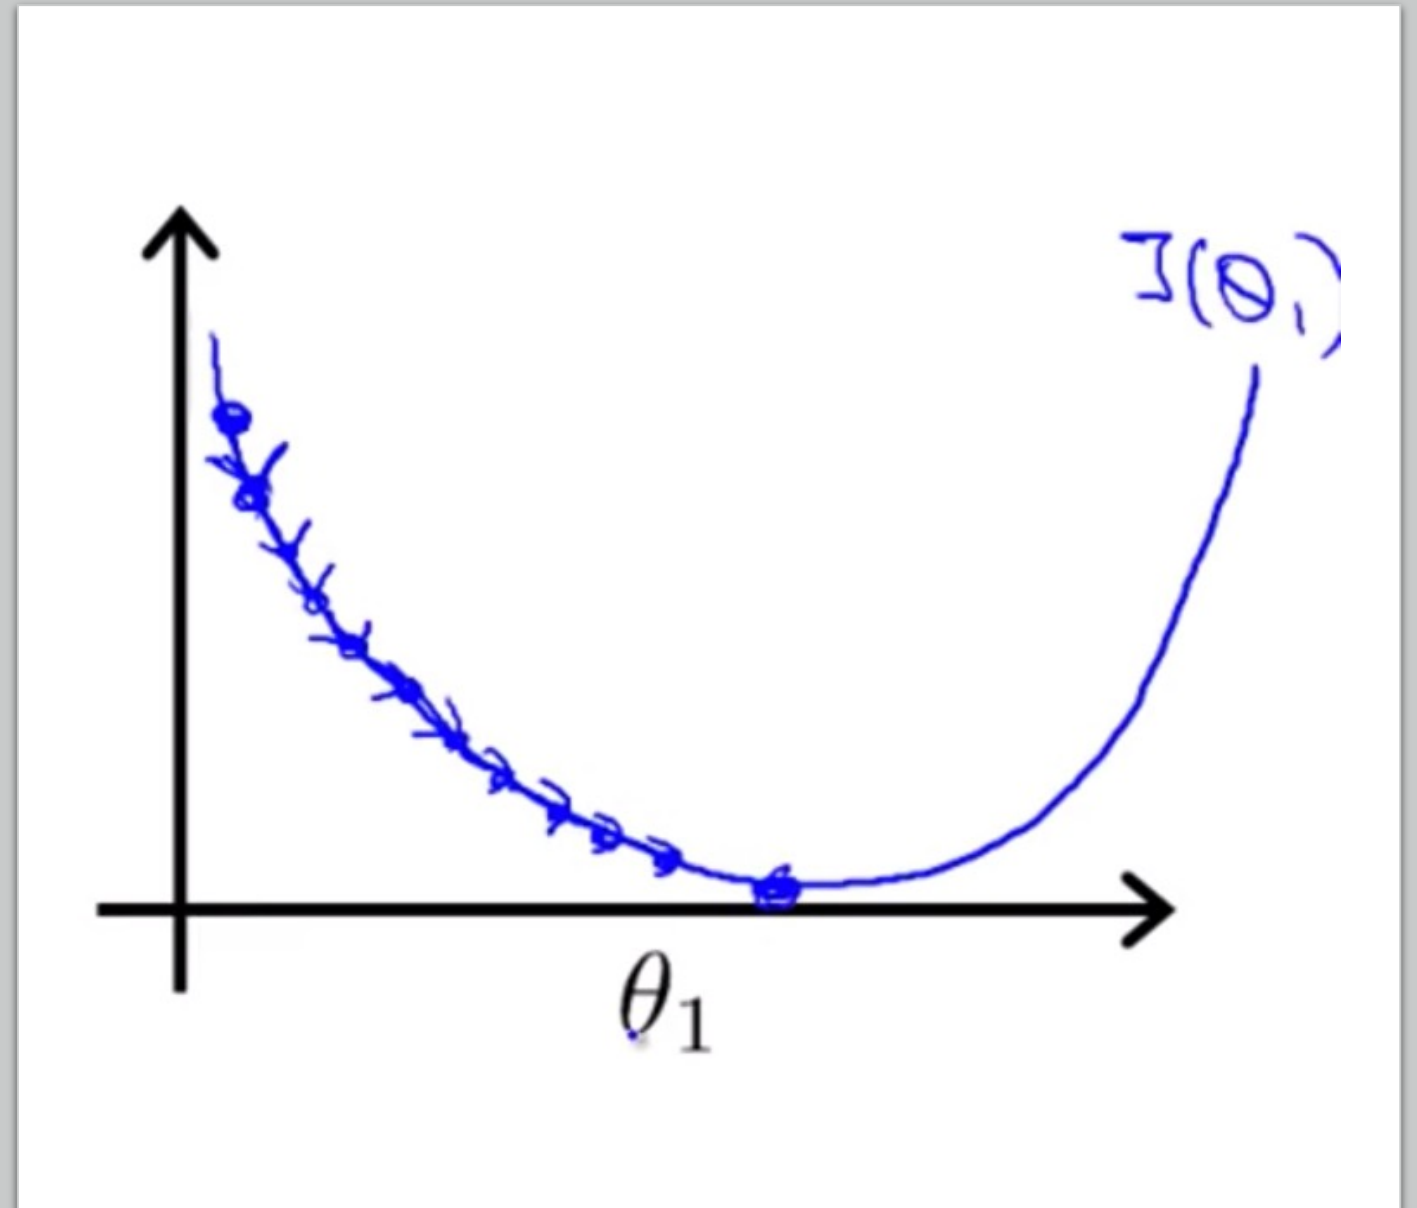
Okay so now what about α ?

- Remember the update rule : $\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1)$
- How does α influence the update of our parameter θ_1 ?
- If α is too small :



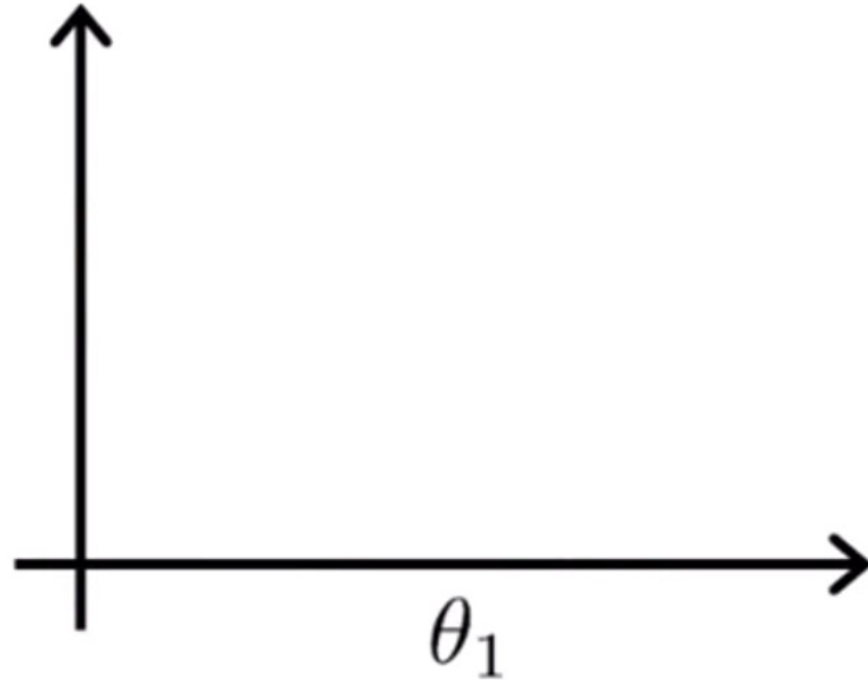
If α is too small

- Many small steps will be taken, which makes Gradient Descent very slow

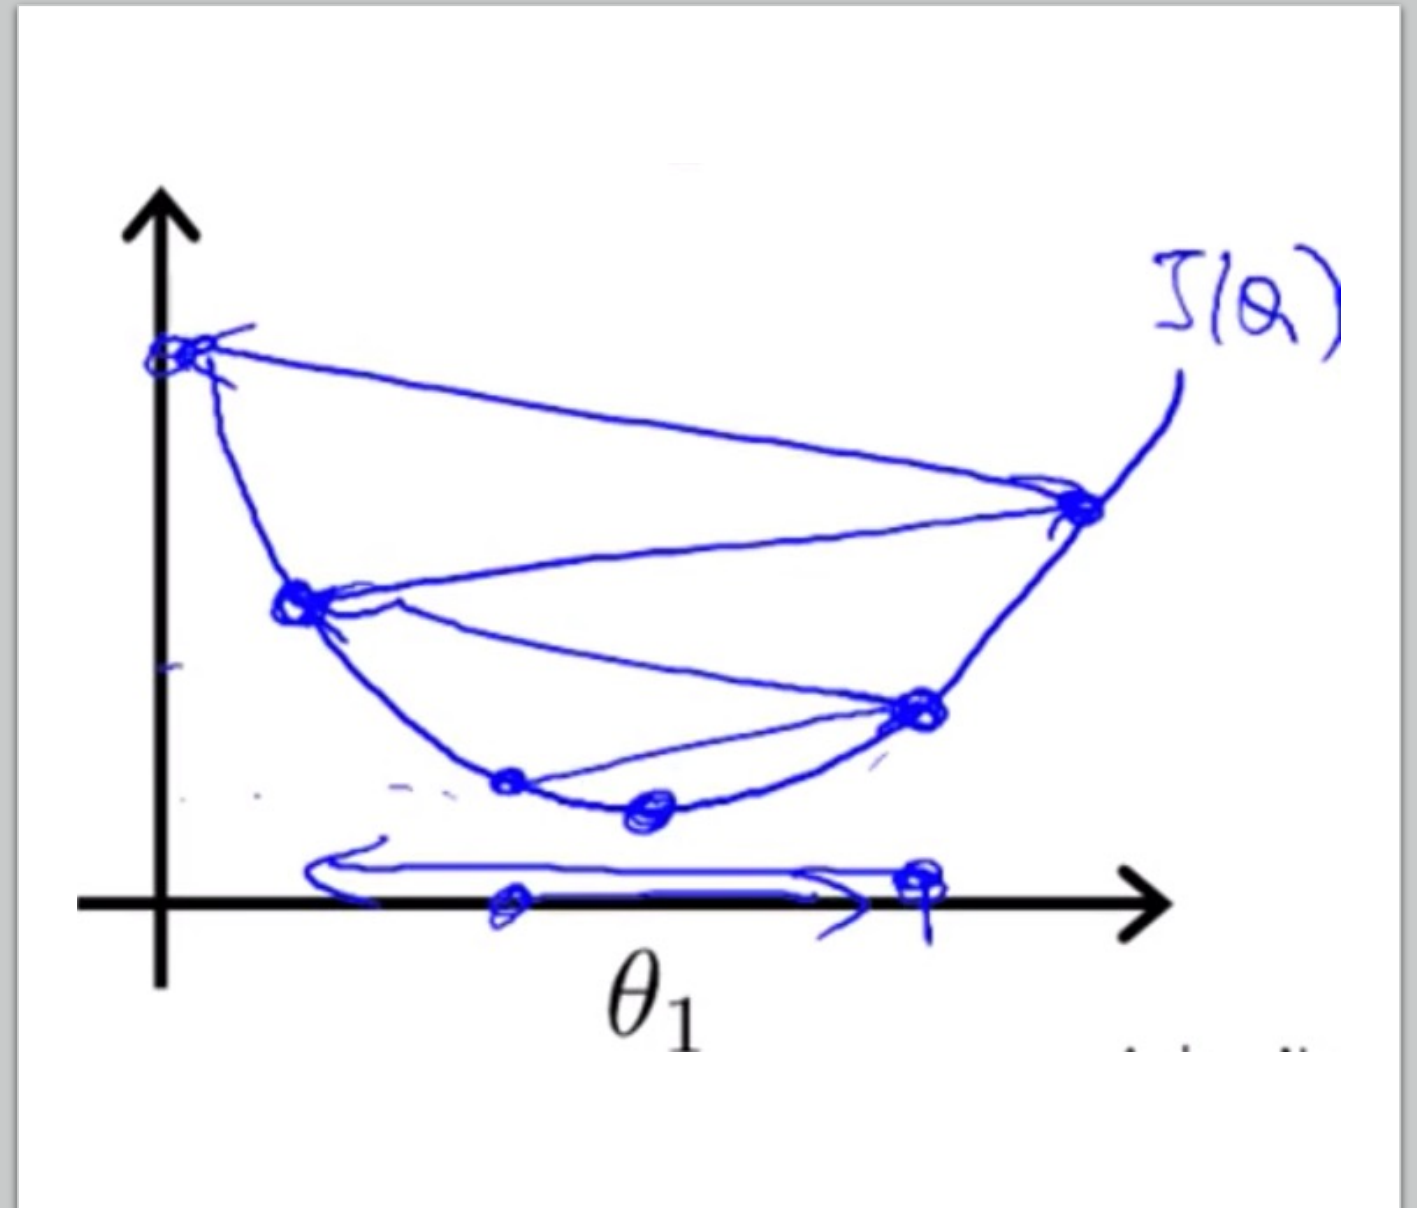


If α is too large...

- Gradient descent may « overshoot », go past the minimum. It may even never converge (never find the minimum) and keep jumping around.



If alpha too large

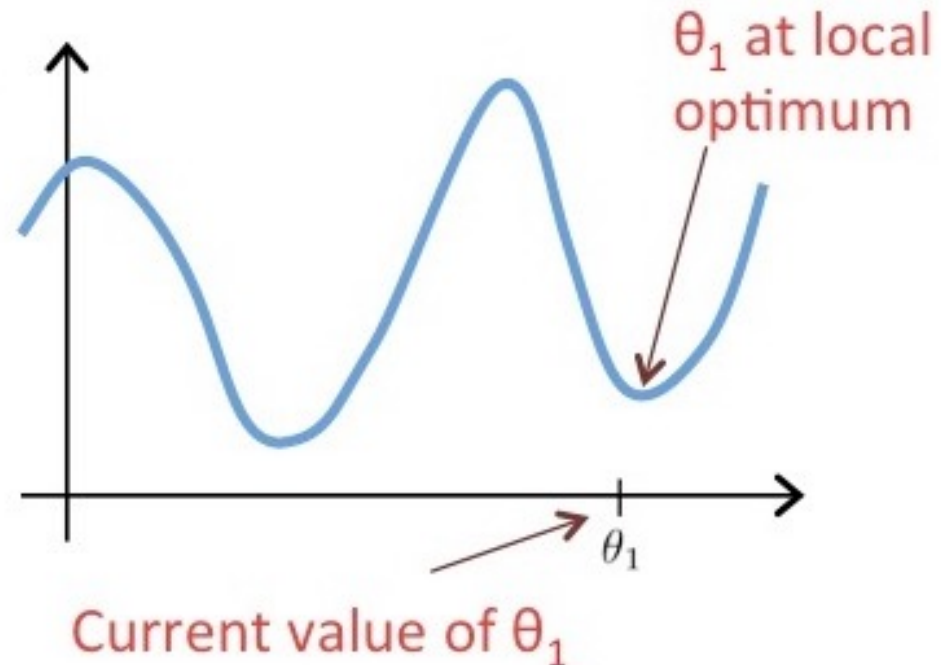


Question

1. Change θ_1 in a random direction ?
2. Move θ_1 in the direction of the global minimum of $J(\theta_1)$?
3. Leave θ_1 unchanged ?
4. Decrease θ_1 ?

Suppose θ_1 is at a local optimum of $J(\theta_1)$, such as shown in the figure.

What will one step of gradient descent $\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1)$ do?



Recap

- To update our parameter with the Gradient Descent algorithm, we perform 2 essential steps :
 1. Compute the derivative of the parameter with respect to the value we want to minimize (ie. our cost: a score to express how good our model is doing)
 2. Take an optimization step/update the parameter. This update will be proportional to the derivative and the learning rate.
- Large derivative (steep tangent line) + large learning rate = big update

Piecing everything together

- This is all we need :
 - A **hypothesis** function (our model)
 - A **cost function** (to tell us how well/bad our model is doing)
 - **Gradient Descent** (to update our parameters and get closer to a better model)

Gradient descent algorithm

repeat until convergence {
 $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$
 (for $j = 1$ and $j = 0$)
}

Linear Regression Model

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Derivatives vs. Partial derivatives

- Except, instead of having a cost function with a single input, we are back to 2 inputs, our 2 parameters θ_0 and θ_1 .
- When we have functions with multiple inputs (known as multivariate functions), computing 1 single derivative is no longer enough!
- The function's « **instantaneous rate of change** » for a given combination of parameters is now determined by 2 values :
 - How does a tiny change in θ_0 change $J(\theta_0, \theta_1)$?
 - How does a tiny change in θ_1 change $J(\theta_0, \theta_1)$?

=> Packed together, these 2 derivatives make up what is referred to as the **gradient**
- Each derivative is a **partial derivative**. (you need both together to get the whole picture !)

Derivatives vs. Partial derivatives

- *Partial Derivative :*

This comes down to calculating the derivative at each input value, treating the other input as a constant

- We pretend for a second that the other input value has basically no effect on the function
 - when looking at θ_0 , we treat θ_1 as a constant
 - when looking at θ_1 , we treat θ_0 as a constant

Partial derivatives and graphs intuition

- To help illustrate things and relate them to our simple Gradient Descent intuition:



Gradient Descent

- Each partial derivative tells us how the function behaves (increases/decreases, quickly/slowly, stays constant...) with respect to a single input
- We can then use this information to know if we should increase or decrease each input to get closer to our minimum cost value !
- **Gradient** : the partial derivatives packed together in a vector
- **Descent** : we want to find the cost function's minimum, using the gradient as a source of information to tell us if the cost is increasing/decreasing with respect to each input.

Update rule

- So we need to figure out the partial derivatives for each parameter !
- the partial derivative of $J(\theta_0, \theta_1)$ with respect to θ_1
- the partial derivative of $J(\theta_0, \theta_1)$ with respect to θ_2

Partial derivatives of $J(\theta_1, \theta_2)$

- You can treat these results as being **given**, in order not to go into the details of the derivation, which involves a couple of the rules mentioned in the optional slides at the end

- General formula

$$\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) = \frac{\partial}{\partial \theta_j} \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i)^2$$

$$\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) = \frac{\partial}{\partial \theta_j} \frac{1}{2m} \sum_{i=1}^m (\theta_0 + \theta_1 x^i - y^i)^2$$

See here for the [MSE derivative](#)

Partial derivatives of $J(\theta_1, \theta_2)$

- Here are the partial derivatives obtained (take these at face value for now):

$$j = 0 : \quad \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i)$$

$$j = 1 : \quad \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i) x^i$$

- These formulas allow us to compute the partial derivatives for each of the parameters, which we can then plug into our Gradient Descent algorithm.

Gradient Descent

- We now have formulas to update our parameters !

$$\begin{aligned} &\text{repeat until convergence} \{ \\ &\quad \theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \\ &\quad \theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x^{(i)} \\ &\} \end{aligned}$$

Quick recap to put things into perspective

- We have :
- a **model**, which is a line :

$$h(x) = \theta_0 + \theta_1 x$$

- a **cost function**, to tell us how good/bad our model is fitting the data:

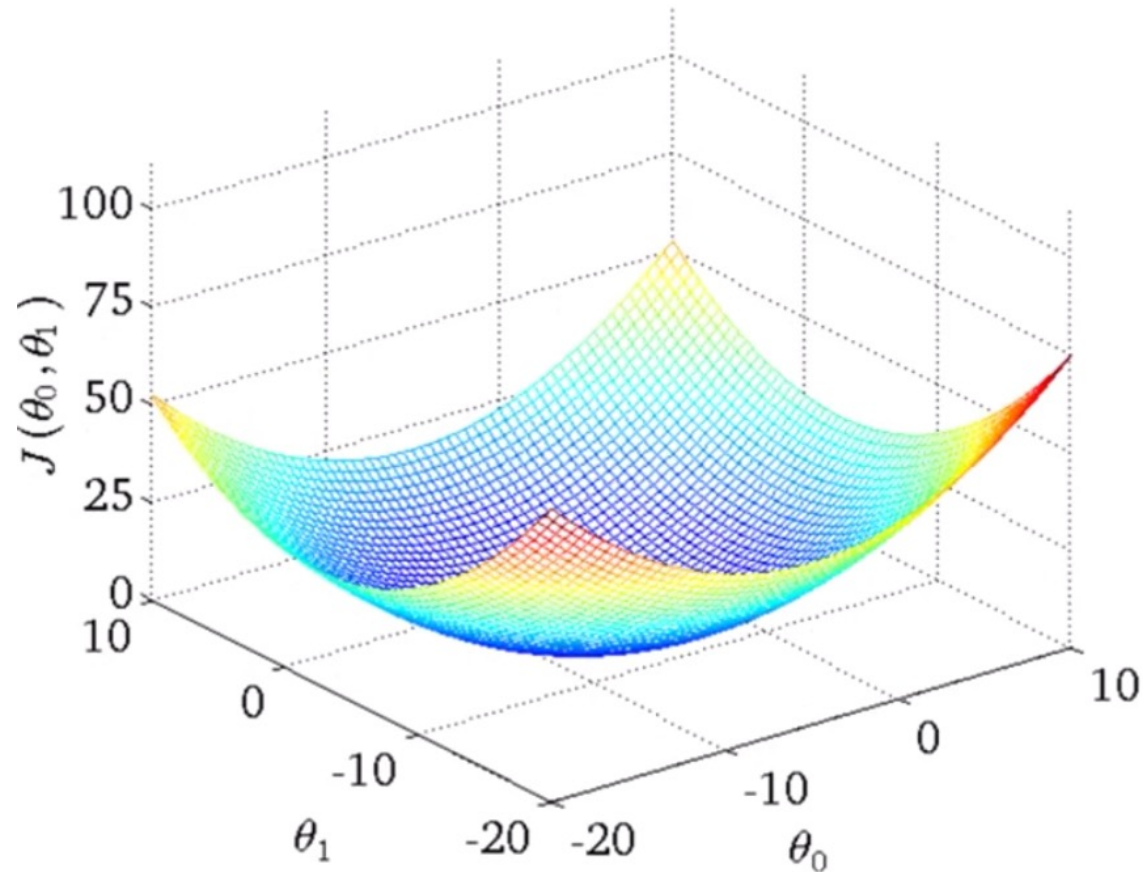
$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$

- **Gradient Descent**, a method to update our parameters so as to minimize the cost function:

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$$

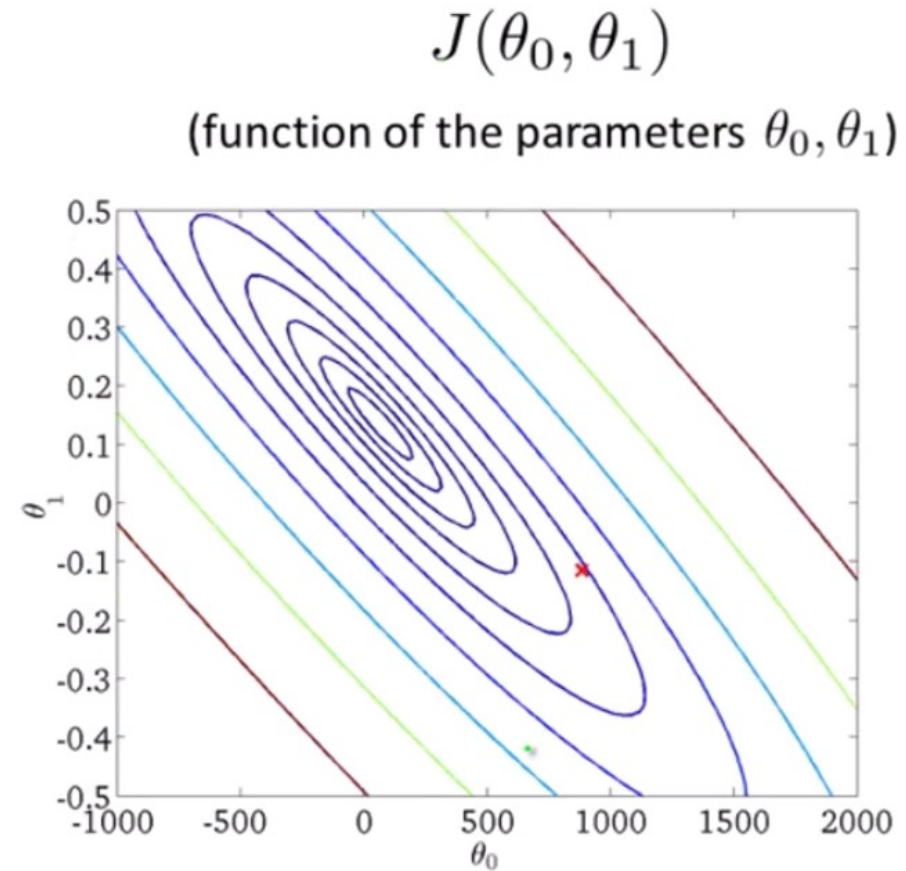
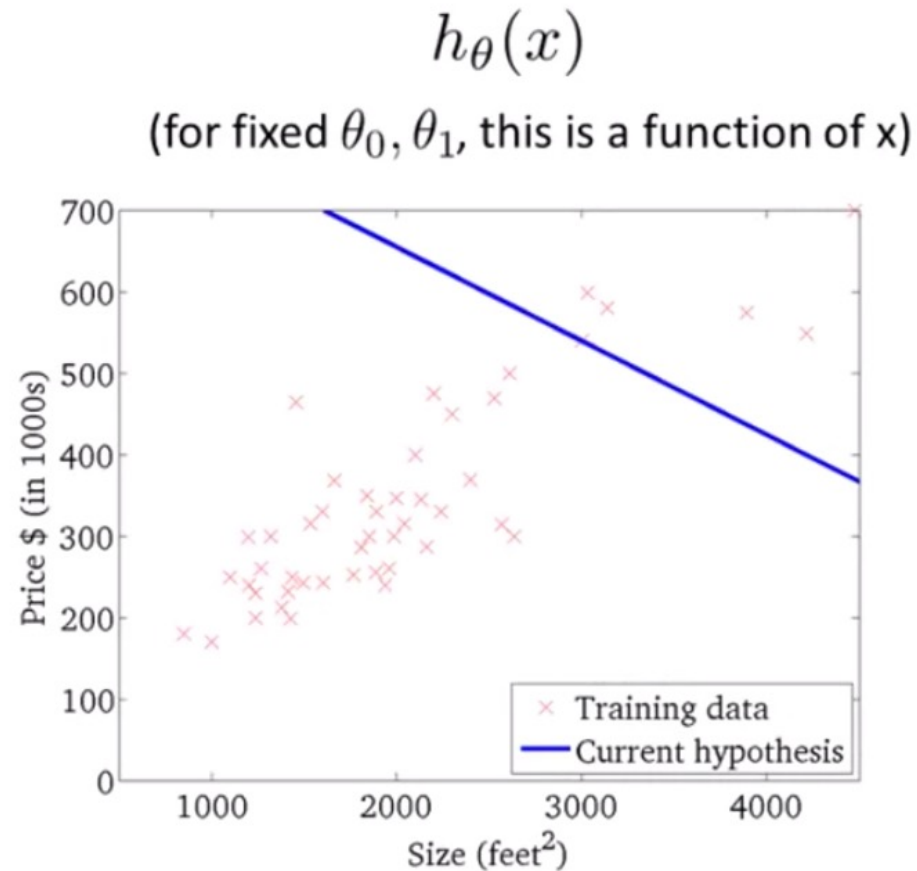
Update examples

- For linear regression, the cost function will always be bowl-shaped



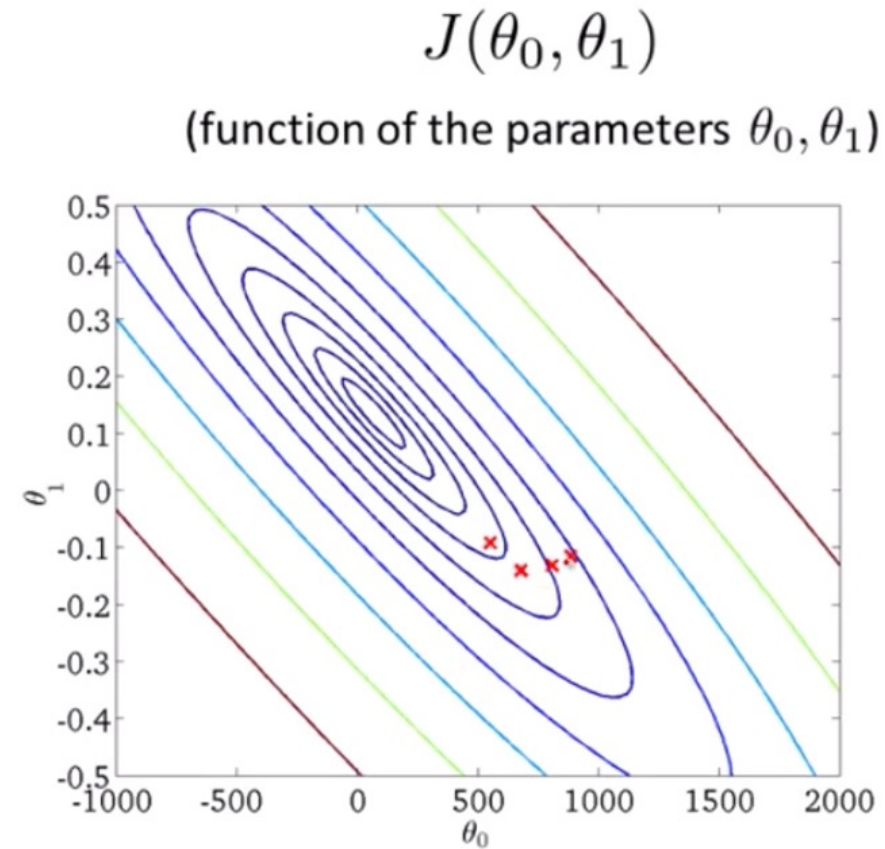
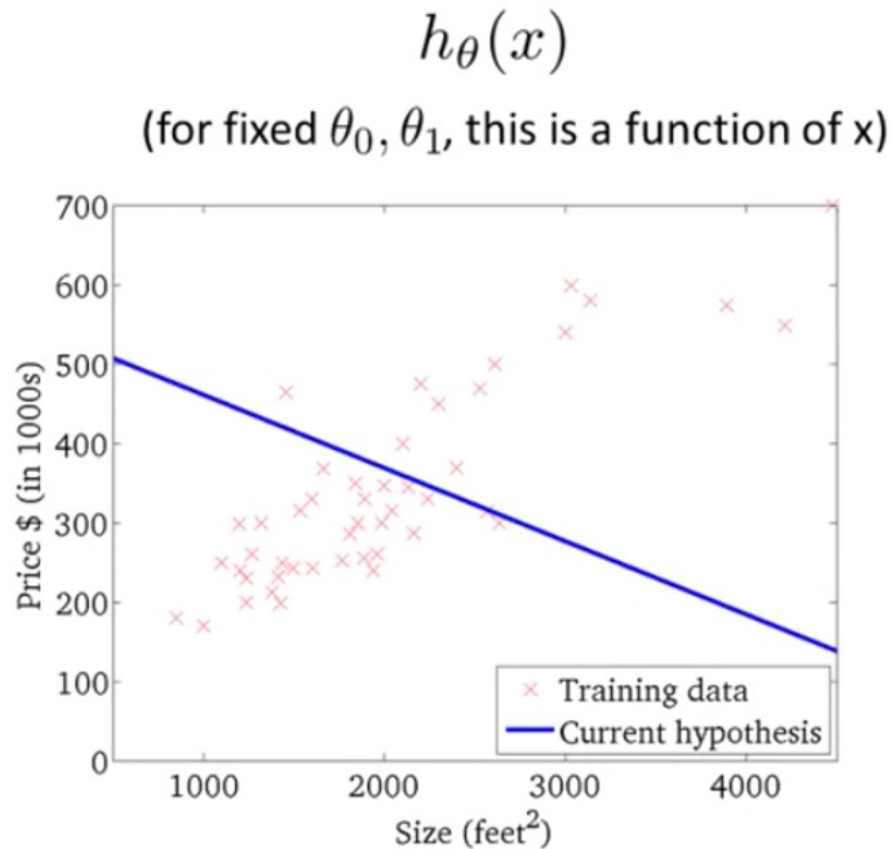
Update examples

Say we initialize our parameters randomly, this is the model and cost :



Update examples

As we take Gradient Descent steps, the model (line) seems to be fitting the data better

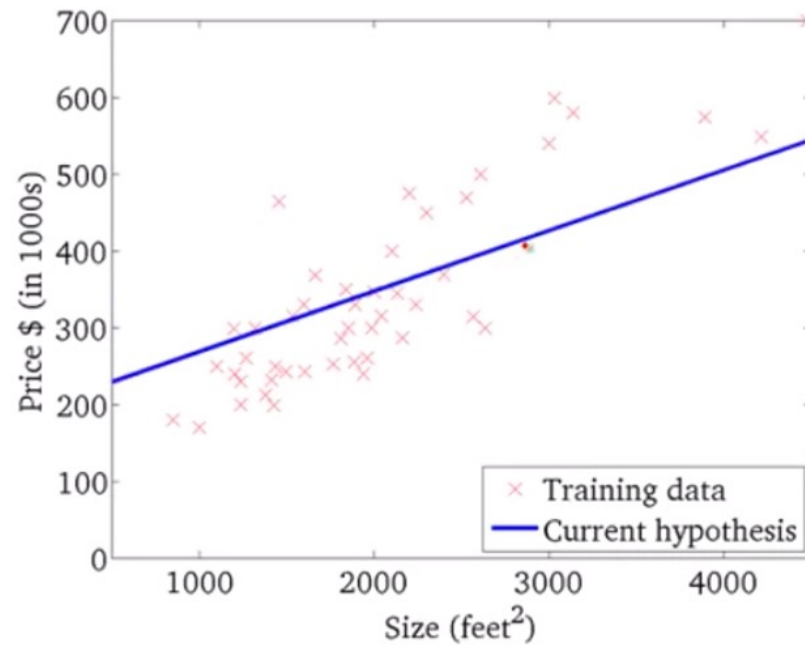


Update examples

Until we reach the global minimum

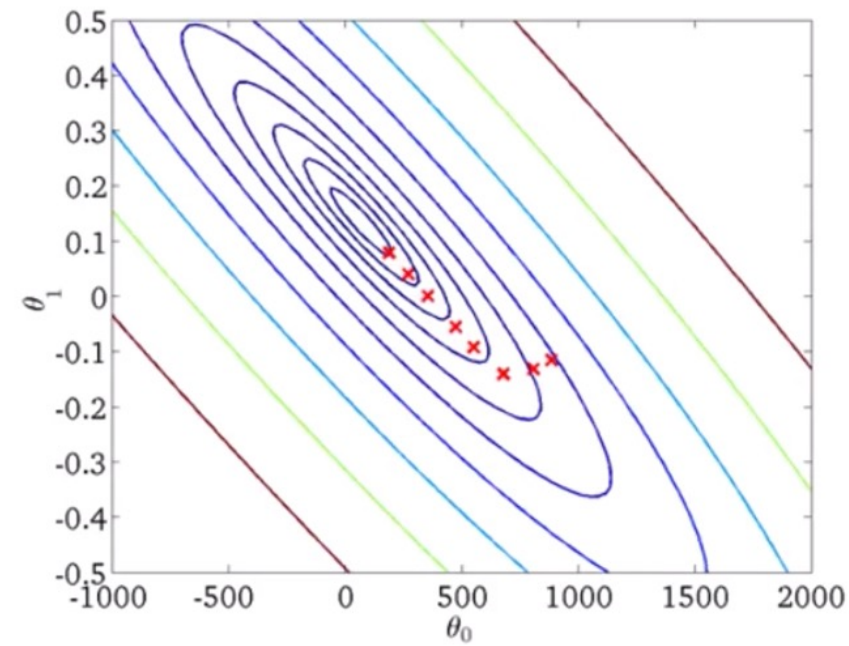
$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)

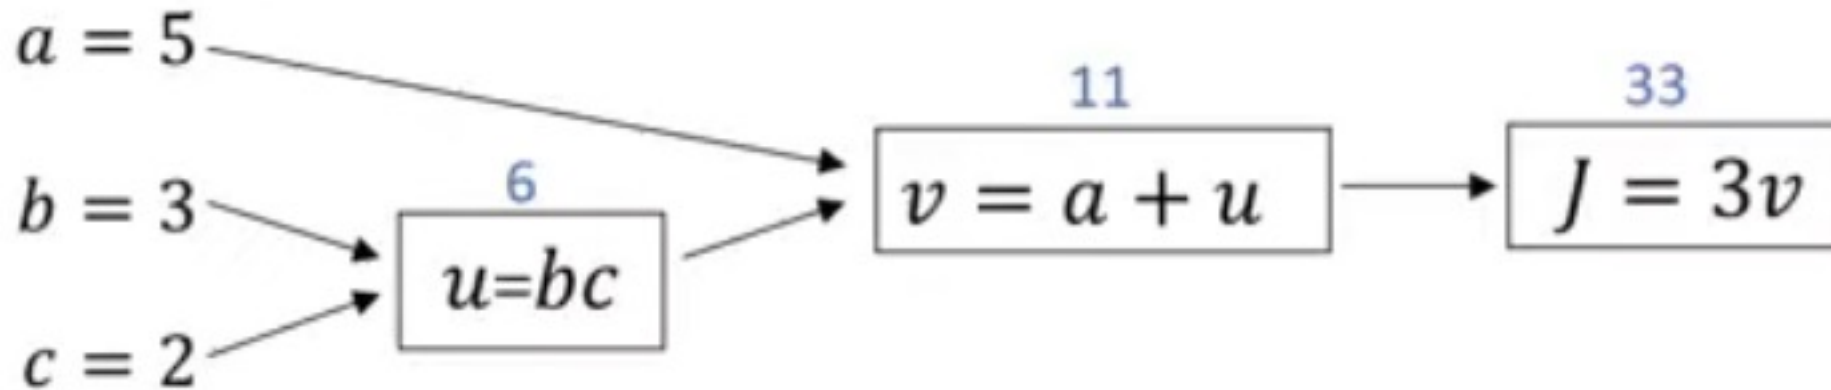


Chain Rule (plus intro to backpropagation)

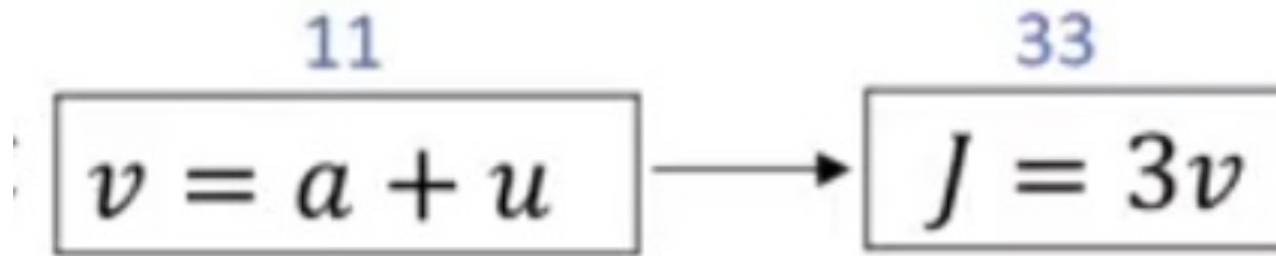
- Just a little intuition on how the derivative of the MSE was calculated.
- The chain rule is the main mathematical tool used to compute partial derivatives and is essential to any algorithm that uses gradient descent !

Chain Rule

- Let's apply derivatives to a computation graph that has multiple nodes, each representing a function :



Finding the derivative the final output J with respect to each node

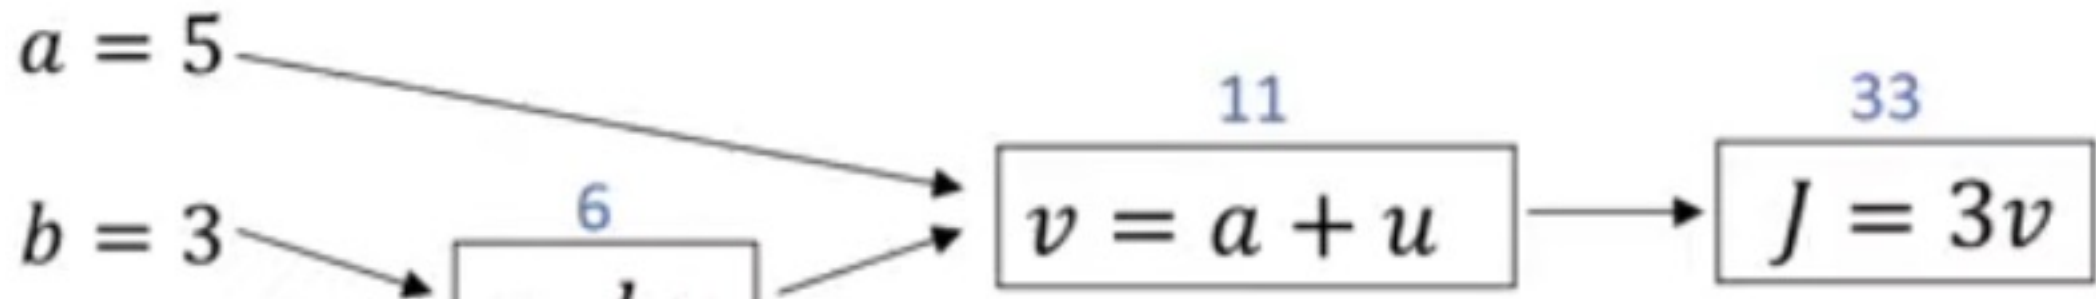


- $\frac{dJ}{dv} = ?$

$$\begin{aligned} v = 11 &\rightarrow 11.001 \\ J = 33 &\rightarrow 33.003 \end{aligned}$$

- $\frac{dJ}{da} = ?$
- To figure this out, let's look at how a affects v , and then how this change in v affects J :

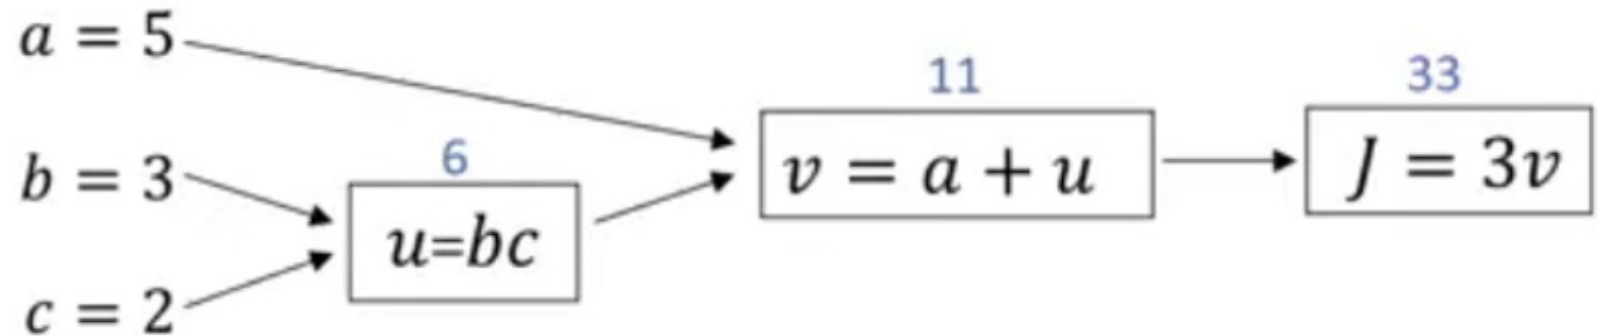
$$\begin{aligned} a = 5 &\rightarrow 5.001 \\ v = 11 &\rightarrow 11.001 \\ J = 33 &\rightarrow 33.003 \end{aligned}$$



- Reasoning by looking at the waterfall effect of nudging a is the basis of the chain rule : $a \rightarrow v \rightarrow J$
- We can know how a nudge in a changes J by multiplying how much this nudge **first** changes v **and then** J .

$$\frac{dJ}{da} = \frac{dJ}{dv} \times \frac{dv}{da}$$

- $\frac{dJ}{du} = ?$
- $\frac{dJ}{db} = ?$
- $\frac{db}{dc} = ?$



- To figure out $\frac{dJ}{du}$, the chain rule tells us

$$\frac{dJ}{du} = \frac{dJ}{dv} \times \frac{dv}{du} = 3 \times 1 = 3$$
- We can re-use the previous computations as we go back through the graph!
- Same applies for the other derivatives.

Multivariate Linear Regression

But what if we have **several features** vs. only 1 ?

| Size (feet ²) | Number of bedrooms | Number of floors | Age of home (years) | Price (\$1000) |
|---------------------------|--------------------|------------------|---------------------|----------------|
| 2104 | 5 | 1 | 45 | 460 |
| 1416 | 3 | 2 | 40 | 232 |
| 1534 | 3 | 2 | 30 | 315 |
| 852 | 2 | 1 | 36 | 178 |
| ... | ... | ... | ... | ... |

Notation

- Features will be denoted by

$$x_1, x_2, \dots, x_n,$$

Where :

Notation:

- n = number of features

$x^{(i)}$ = input (features) of i^{th} training example.

$x_j^{(i)}$ = value of feature j in i^{th} training example.

Notation

- With this notation, the example $\mathbf{x}^{(2)}$ is a 4-D vector :

- $\mathbf{x}^{(2)} = \begin{bmatrix} 1416 \\ 3 \\ 2 \\ 40 \end{bmatrix}$

- $x_3^{(2)} = 2$

| x_1 Size (feet ²) | x_2 Number of bedrooms | x_3 Number of floors | x_4 Age of home (years) | Price (\$1000) |
|------------------------------------|--------------------------------|------------------------------|---------------------------------|----------------|
| 2104 | 5 | 1 | 45 | 460 |
| 1416 | 3 | 2 | 40 | 232 |
| 1534 | 3 | 2 | 30 | 315 |
| 852 | 2 | 1 | 36 | 178 |
| ... | ... | ... | ... | ... |

The model

- Previously $h_{\theta}(x) = \theta_0 + \theta_1 x$

- Now,

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

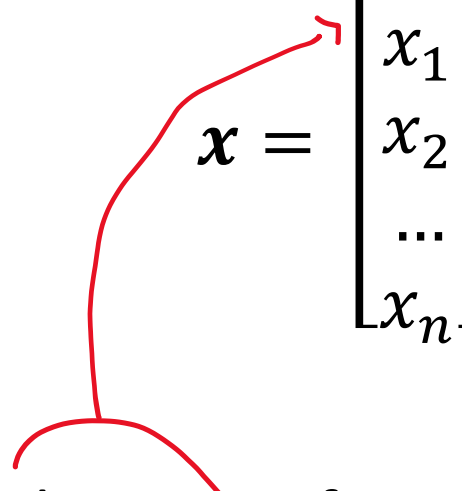
- We can no longer represent the model using a 2D graph...

Gathering features and parameters into vectors

- For convenience, let's group up features and parameters into vectors:

$$\bullet \mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \dots \\ x_n \end{bmatrix} \quad \mathbf{\theta} = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \dots \\ \theta_n \end{bmatrix}$$

Size => n ~~≠~~ Size => n+1


$$\mathbf{x} = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \dots \\ x_n \end{bmatrix}$$

- To avoid this size mismatch we can add a « dummy » feature, $x_0 = 1$

The model formula simplified

$$h_{\theta}(x) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$$

Instead of using a loop, which is slow, we can vectorize both the inputs and the parameters and compute their dot product !

- $h_{\theta}(x) = \boldsymbol{\theta}^T \boldsymbol{x}$

- $[\theta_0 \ \theta_1 \ \dots \ \theta_n] \begin{bmatrix} x_0 \\ x_1 \\ \dots \\ x_n \end{bmatrix}$

Multivariate Gradient Descent

- The intuitions and formulas we saw previously are the same, there are just more partial derivatives to compute !

New algorithm ($n \geq 1$):

Repeat {

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

(simultaneously update θ_j for
 $j = 0, \dots, n$)

}

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_1^{(i)}$$

$$\theta_2 := \theta_2 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_2^{(i)}$$

...

Derivative rules (optional)

- Just so you are aware, no need to learn these by heart. But useful if you want to try and derive a function on your own !
- You can find these kinds of « cheatsheets » online if you need to.

| | Function $f(x)$ | Derivative with respect to x |
|----|-----------------|--------------------------------|
| 1 | a | 0 |
| 2 | x | 1 |
| 3 | ax | a |
| 4 | x^2 | $2x$ |
| 5 | x^a | ax^{a-1} |
| 6 | a^x | $\log(a)a^x$ |
| 7 | $\log(x)$ | $1/x$ |
| 8 | $\log_a(x)$ | $1/(x \log(a))$ |
| 9 | $\sin(x)$ | $\cos(x)$ |
| 10 | $\cos(x)$ | $-\sin(x)$ |
| 11 | $\tan(x)$ | $\sec^2(x)$ |

Derivative rules (optional)

- More useful rules

| | Function | Derivative |
|-----------------|---------------|------------------------------------|
| Sum Rule | $f(x) + g(x)$ | $f'(x) + g'(x)$ |
| Difference Rule | $f(x) - g(x)$ | $f'(x) - g'(x)$ |
| Product Rule | $f(x)g(x)$ | $f'(x)g(x) + f(x)g'(x)$ |
| Quotient Rule | $f(x)/g(x)$ | $[g(x)f'(x) - f(x)g'(x)]/[g(x)]^2$ |
| Reciprocal Rule | $1/f(x)$ | $-[f'(x)]/[f(x)]^2$ |
| Chain Rule | $f(g(x))$ | $f'(g(x))g'(x)$ |

Quick but useful example using the rules (optional)

- $f(\theta_1) = \frac{(\theta_1 x_1 - y_1)^2}{2}$
- Power rule : $x^2 \Rightarrow 2x$
- Scalar multiplication rule: $ax \Rightarrow a$
- Chain rule: $\frac{d}{dx} f(g(h(x))) \Rightarrow \frac{df}{dg} \times \frac{dg}{dh} \times \frac{dh}{dx}$

(Optional)

- $f(a) = \frac{(ax - y)^2}{2}$
- Let's decompose this into 3 functions:
 - $g(a) = ax - y$
 - $h(g) = g^2$ where g is $(ax - y)$
 - $i(h) = \frac{h}{2}$ where h is $(ax - y)^2$
- Let's derive these functions 1 by 1 using the rules in the previous slides.

(Optional)

- $g'(a) = \frac{d}{da}(ax - y) = x$
 - $h'(g) = \frac{d}{dg}(g^2) = 2g$
 - $i'(h) = \frac{d}{dh}\left(\frac{1}{2}h\right) = \frac{1}{2}$
- Using the chaine rule, we multiply these dervatives to get the derivative of our original function :

$$\begin{aligned}f'(a) &= \frac{di}{dh} \times \frac{dh}{dg} \times \frac{dg}{da} \\&= \frac{1}{2} \times 2(ax - y) \times x \\&= (ax - y) \times x\end{aligned}$$