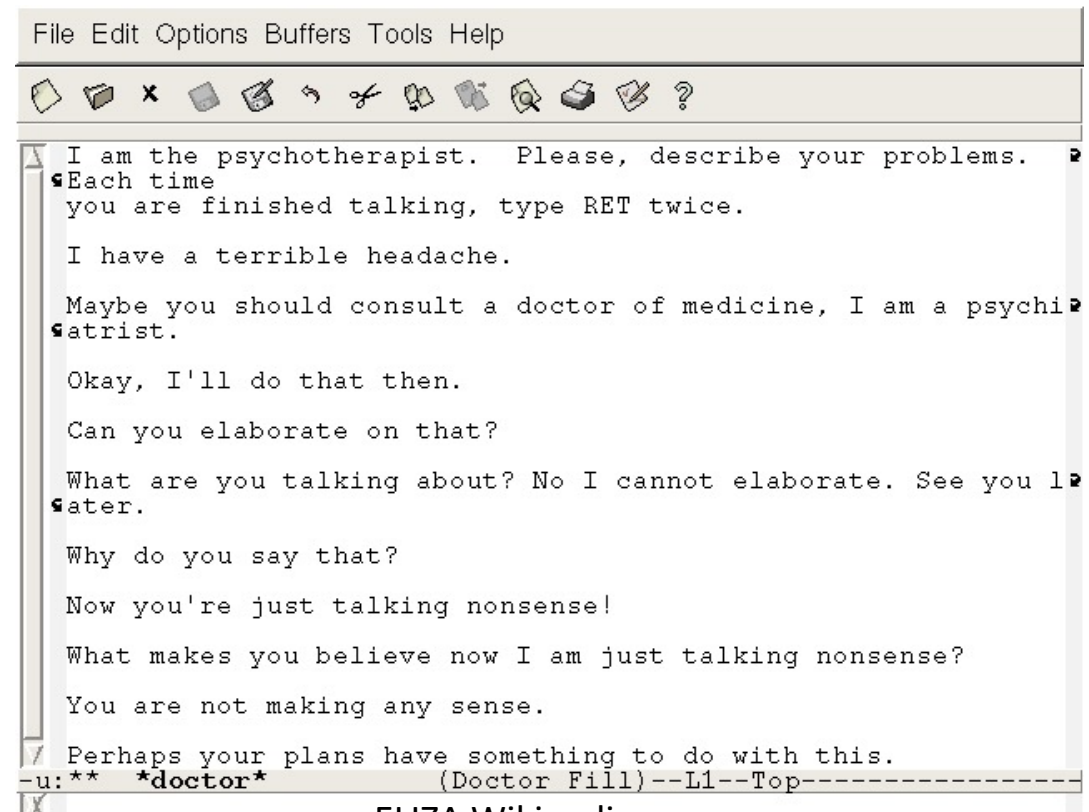


# Perceptron and Multi-Layer Perceptron (MLP)

# History of ANNs

- Interest in ANNs came in waves :
  - 1960s : people thought we would soon be conversing with truly intelligent machines.
    - ELIZA, Weisenbaum 1964 :
  - When this wasn't the case, funding went elsewhere...



```
File Edit Options Buffers Tools Help
[Icons]
I am the psychotherapist. Please, describe your problems.
Each time
you are finished talking, type RET twice.

I have a terrible headache.

Maybe you should consult a doctor of medicine, I am a psychiatrist.

Okay, I'll do that then.

Can you elaborate on that?

What are you talking about? No I cannot elaborate. See you later.

Why do you say that?

Now you're just talking nonsense!

What makes you believe now I am just talking nonsense?

You are not making any sense.

Perhaps your plans have something to do with this.
-u: ** *doctor* (Doctor Fill)--L1--Top-----
```

ELIZA Wikipedia page

# History of ANNs


- 1980s : revival of interest, new architectures invented, better training techniques, but progress was slow...
- by the 90s, other machine learning techniques were invented (SVMs) => seemed to offer better results and stronger theoretical foundations.

# History of ANNs

- Present day : another wave of interest, but this one seems to be different :
  - Now a huge quantity of data available to train neural networks, and ANNs frequently outperform other ML techniques on very large and complex problems.
  - Tremendous increase in computing power since the 1990s now makes it possible to train large neural networks in a « reasonable » amount of time.
  - Training algorithms have been improved.
  - Possible to create ANNs on a massive scale (100 Trillion params in GPT4)
  - ANNs seem to have entered a virtuous circle of funding and progress.


# Biological Neurons

- Biological neurons receive short electrical impulses called signals from other neurons via synapses.
- When a neuron receives a sufficient number of signals from other neurons, it fires its own signals.



# Biological Neurons

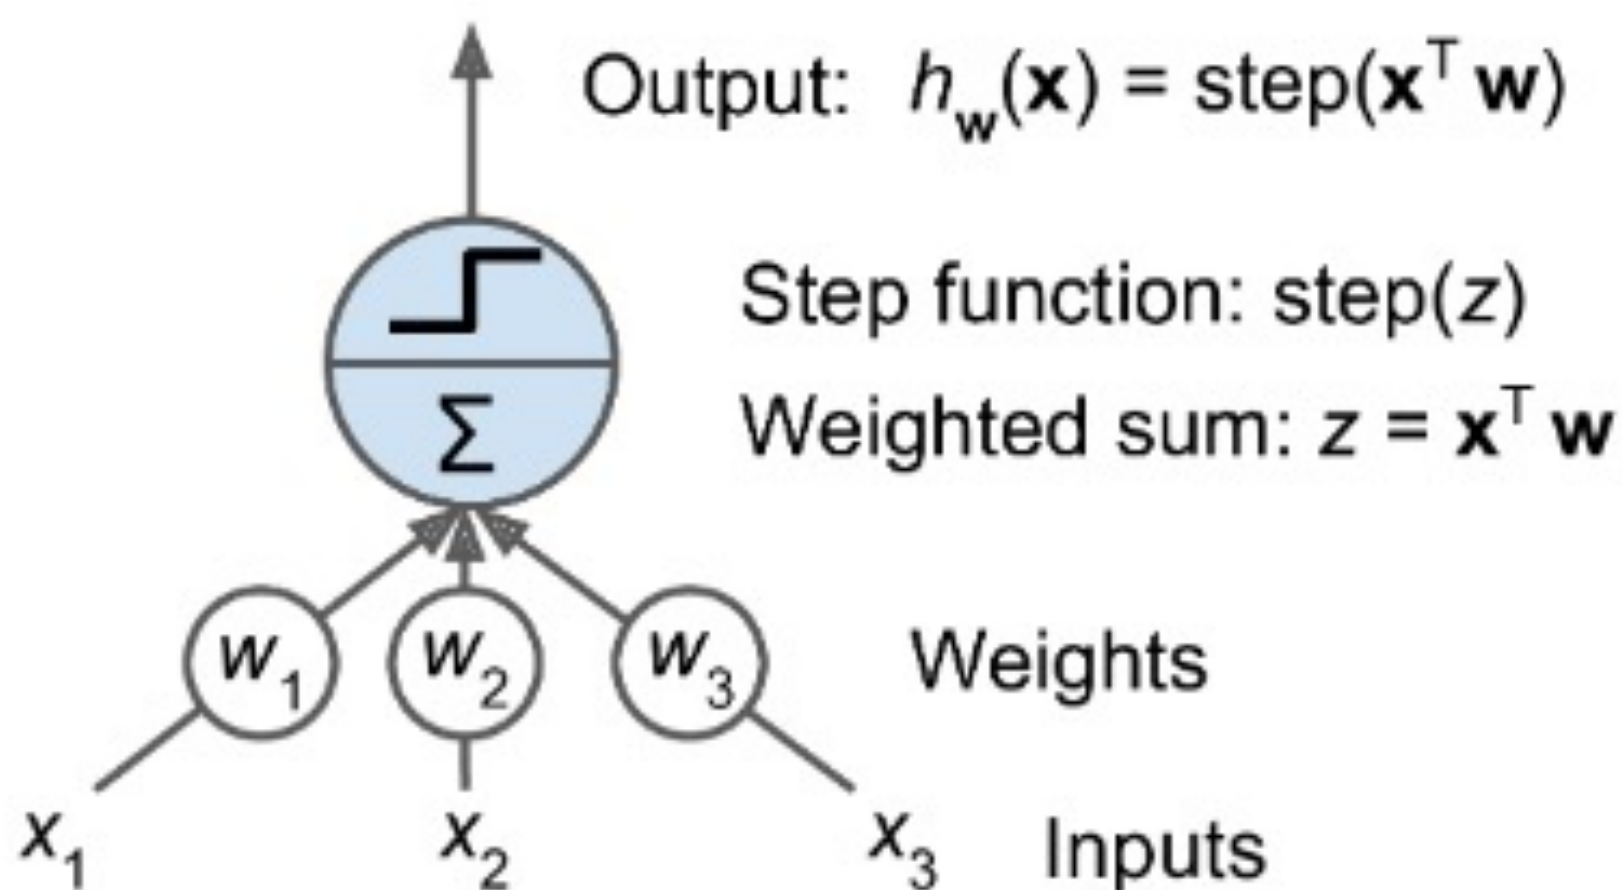
- Organized in a vast network of billions of neurons, each neuron is typically connected to thousands of other neurons.
- Research suggests that neurons are often organized in consecutive layers



Drawing of a cortical lamination by S. Ramon y Cajal (public domain)

# The Perceptron

- One of the simplest ANN architectures, invented in 1957 by Frank Rosenblatt
- These artificial neurons are called *Threshold Logic Units*. We will see why.





# The Step Function

- Most commonly used is the *Heaviside Step Function* :

$$\text{heaviside}(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$$

- Does this unit ring any bells...?


# Binary Classification

- A single TLU can be used for simple linear binary classification.
- It computes a linear combination of the inputs and if the result exceeds a threshold, it outputs the positive class or else outputs the negative class.

# Perceptron

- The Perceptron can refer to this single TLU, but can also refer to a single *layer* of TLUs.
- When all the neurons in a layer are connected to every neuron in the previous layer, it is called a *fully connected layer* or a *dense layer*.

# Perceptron Diagram



# Perceptron

- Common to draw special passthrough neurons called *input neurons*
- Also, an extra bias feature is can be added (our dummy feature from linear/logistic regression), but this can be handled differently
- This particular Perceptron setup can classify instances simultaneously into three different binary classes, which makes it a multi-output classifier

# Computing the outputs of a fully connected layer

$$h_{\mathbf{W},\mathbf{b}}(\mathbf{X}) = \textit{activation}(\mathbf{X}\mathbf{W} + \mathbf{b})$$

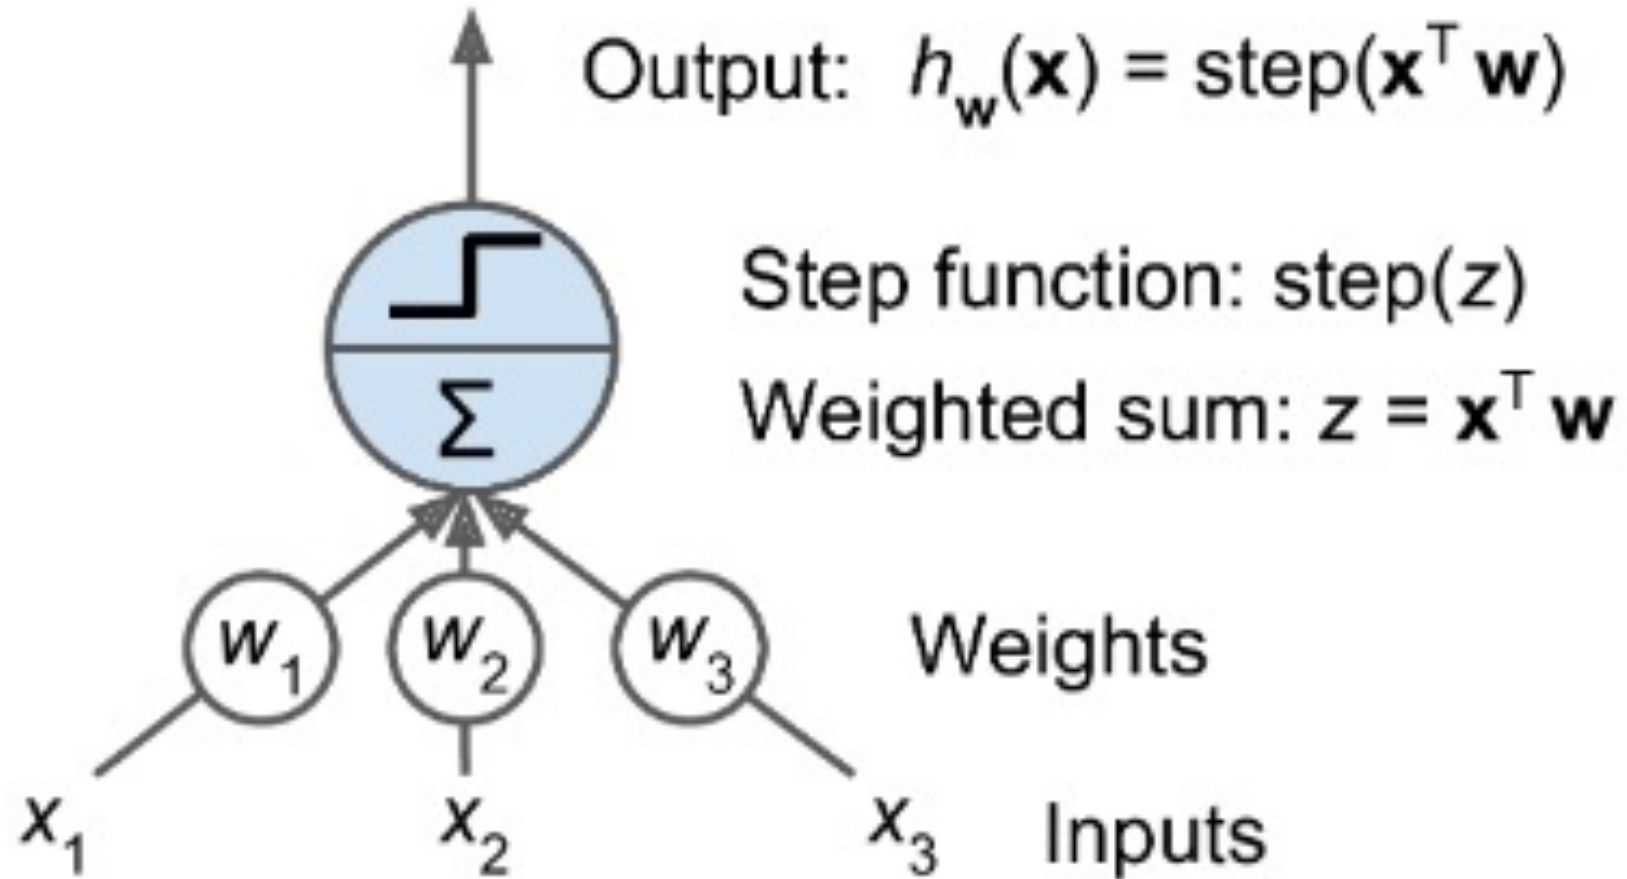
(no dummy features need to be added to  $\mathbf{X}$  in this case)

- $\mathbf{X}$  represents the matrix of input features. It has one row per instance, one column per feature.
- The weight matrix  $\mathbf{W}$  contains all the connection weights except for the ones from the bias neuron. It has one row per input neuron and one column per artificial neuron in the layer.
- The bias vector  $\mathbf{b}$  contains all the connection weights between the bias neuron and the artificial neurons. It has one bias term per artificial neuron.
- The activation function: when the artificial neurons are TLUs, it is a step function. We will discuss other activation functions.

# How is the Perceptron trained ?

- In his book *The Organization of Behavior*, published in 1949, Donald Hebb suggested that *when a biological neuron often triggers another neuron, the connection between these two neurons grows stronger.*
- **“Cells that fire together, wire together.”** (Siegfried Löwel)
- Perceptrons are trained using a variant of this rule that takes into account the error made by the network
- For every output neuron that produced a wrong prediction, it *reinforces the connection weights from the inputs that would have contributed to the correct prediction.*
- I.e. Increases the weight in the connection

## Example





# What can the Perceptron Learn ?

- Cannot learn complex patterns in the data
- However, if the training instances are linearly separable, Rosenblatt demonstrated that this algorithm would converge to a solution.

# Sklearn Perceptron

- Scikit-Learn provides a Perceptron class that implements a **single TLU** network (on the iris dataset example here):

```
import numpy as np
from sklearn.datasets import load_iris
from sklearn.linear_model import Perceptron


iris = load_iris()
X = iris.data[:, (2, 3)] # petal length, petal width
y = (iris.target == 0).astype(np.int) # Iris Setosa?

per_clf = Perceptron()
per_clf.fit(X, y)

y_pred = per_clf.predict([[2, 0.5]])
```


# The Multi-Layer Perceptron: Why Multiple Layers ?

- Certain trivial problems prove unsolvable for the simple perceptron.
- The XOR (*exclusive OR*) problem in particular is one of the most famous examples.




# Why Multiple Layers ?

- Turns out some of these limitations can be eliminated by **adding layers** :
- This network solves the XOR problem for example
- With inputs (0,0) or (1, 1) the network outputs 0, and with inputs (0, 1) or (1, 0) it outputs 1.
- All connections have a weight equal to 1, except the four connections where the weight is shown.



# Multi-Layer Perceptron



# Multi-Layer Perceptron


- An MLP is composed of one (passthrough) input layer, one or more layers of TLUs, called *hidden layers*, and one final layer of TLUs called the *output layer*
- The layers close to the input layer are usually called the *lower* layers, and the ones close to the outputs are usually called the *upper* layers.
- Every layer except the output layer includes a *bias neuron* and is *fully connected* to the next layer.
- The signal flows only in one direction (from the inputs to the outputs), so this architecture is an example of a **feedforward neural network** (FNN).

# How do we train this neural network ?

- 1986 : Rumelhart & al. Introduced the ***Backpropagation*** algorithm, still used today.
- Basically a version of Gradient Descent :
  - It's able to compute the gradient of the network's error with regards to *every single model parameter*. This is done using the chain rule.
  - In other words it finds out how each connection weight and each bias term should be tweaked in order to reduce the error.
  - Then a gradient descent step is performed (e.g. for the 1st neuron of the 1st layer) :
$$weights_{N=1,L=1} = weights_{N=1,L=1} - learningRate * \frac{\partial Cost}{\partial weights_{N=1,L=1}}$$

# Chain Rule (slides from previous class on linear regression)


- Let's apply derivatives to a computation graph that has multiple nodes, each representing a function :





- $\frac{dJ}{da} = ?$
- To figure this out, let's look at how  $a$  affects  $v$ , and then how this change in  $v$  affects  $J$  :

$$\begin{aligned} a = 5 &\rightarrow 5.001 \\ v = 11 &\rightarrow 11.001 \\ J = 33 &\rightarrow 33.003 \end{aligned}$$




- Reasoning by looking at the waterfall effect of nudging  $a$  is the basis of the chain rule :  $a \rightarrow v \rightarrow J$
- We can know how a nudge in  $a$  changes  $J$  by multiplying how much this nudge **first** changes  $v$  **and then**  $J$ .

$$\frac{dJ}{da} = \frac{dJ}{dv} \times \frac{dv}{da}$$

# Activation functions


- One key change was made to the neurons for backprop to work:
- The *step* function was replaced with the logistic function.
- Step function contains only flat segments (derivative is 0 everywhere) vs. Sigmoid/logistic function which has a non-zero derivative everywhere
- This is important for the chain rule to work.



Sigmoid


$$f(x) = \begin{cases} 0 & \text{if } 0 > x \\ 1 & \text{if } x \geq 0 \end{cases}$$

Unit step (threshold)



# Other Activation functions


- *The hyperbolic tangent function*  $\tanh(z) = 2\sigma(2z) - 1$   
Just like the logistic function it is S-shaped, continuous, and differentiable, but its output value ranges from  $-1$  to  $1$  (instead of  $0$  to  $1$  in the case of the logistic function)



# Other Activation functions

- The Rectified Linear Unit function:  $\text{ReLU}(z) = \max(0, z)$   
It is continuous but unfortunately not differentiable at  $z = 0$  (the slope changes abruptly, which can make Gradient Descent bounce around), and its derivative is 0 for  $z < 0$ .
- However, in practice it works very well and has the advantage of being fast to compute.

## ReLU Function



# Why do we need activation functions ?

- Why does a neuron not just output the weighted sum of its inputs and pass it on to the next neurons ?
- Why is a **non-linear** activation function needed ?



# Why do we need activation functions ?

- => If you compose several linear transformations, you actually get 1 linear transformation, which doesn't allow us to solve complex problems....
- If  $f(x) = 2x + 3$  and  $g(x) = 5x - 1$ , then composing these two linear functions gives us another linear function:

$$f(g(x)) = 2(5x - 1) + 3 = \mathbf{10x + 1}$$

- So if there isn't a non-linear function between the layers, then even a deep stack of layers is equivalent to a single layer...
- We will see this with a practical example in pytorch later

# Activation Functions and their derivatives





# Extra (illustrated) resources

- Jay Alammar's [blogposts](#) on neural nets
- 3Bue1Brown's [videos](#) on neural nets
- To [play around](#) with a neural net
- [The Absolutely Simplest Backpropagation Example](#)
- [Why neural nets can learn almost anything](#)

# Classification MLPs

- **Binary Classification :**
- 1 single output neuron using the logistic activation function.
- It outputs a number between 0 and 1, which you can interpret as the estimated probability of the positive class.
- The estimated probability of the negative class is equal to  $1 - \text{that number}$ .
- What happens in this output neuron is equivalent to what we saw for logistic regression.
- The only difference is that the input values are now the computed outputs from the neurons in the previous layer vs. the features.


# Classification MLPs

- **Multiclass Classification:**


- If each example fed into the network can only belong to a single class (the classes are mutually exclusive — there can't be a combination of classes, only 1)
- Then you need 1 output neuron *per class* and you need to use the **softmax** activation function for the whole output layer.
- This function ensures the scores produced by each output neuron are between 0 and 1, and add up to 1.

- [Softmax wikipedia](#)
- [Softmax vs argmax](#)

SOFTMAX TRANSFORMS A VECTOR OF NUMBERS INTO A VECTOR OF RELATIVE "PROBABILITIES"



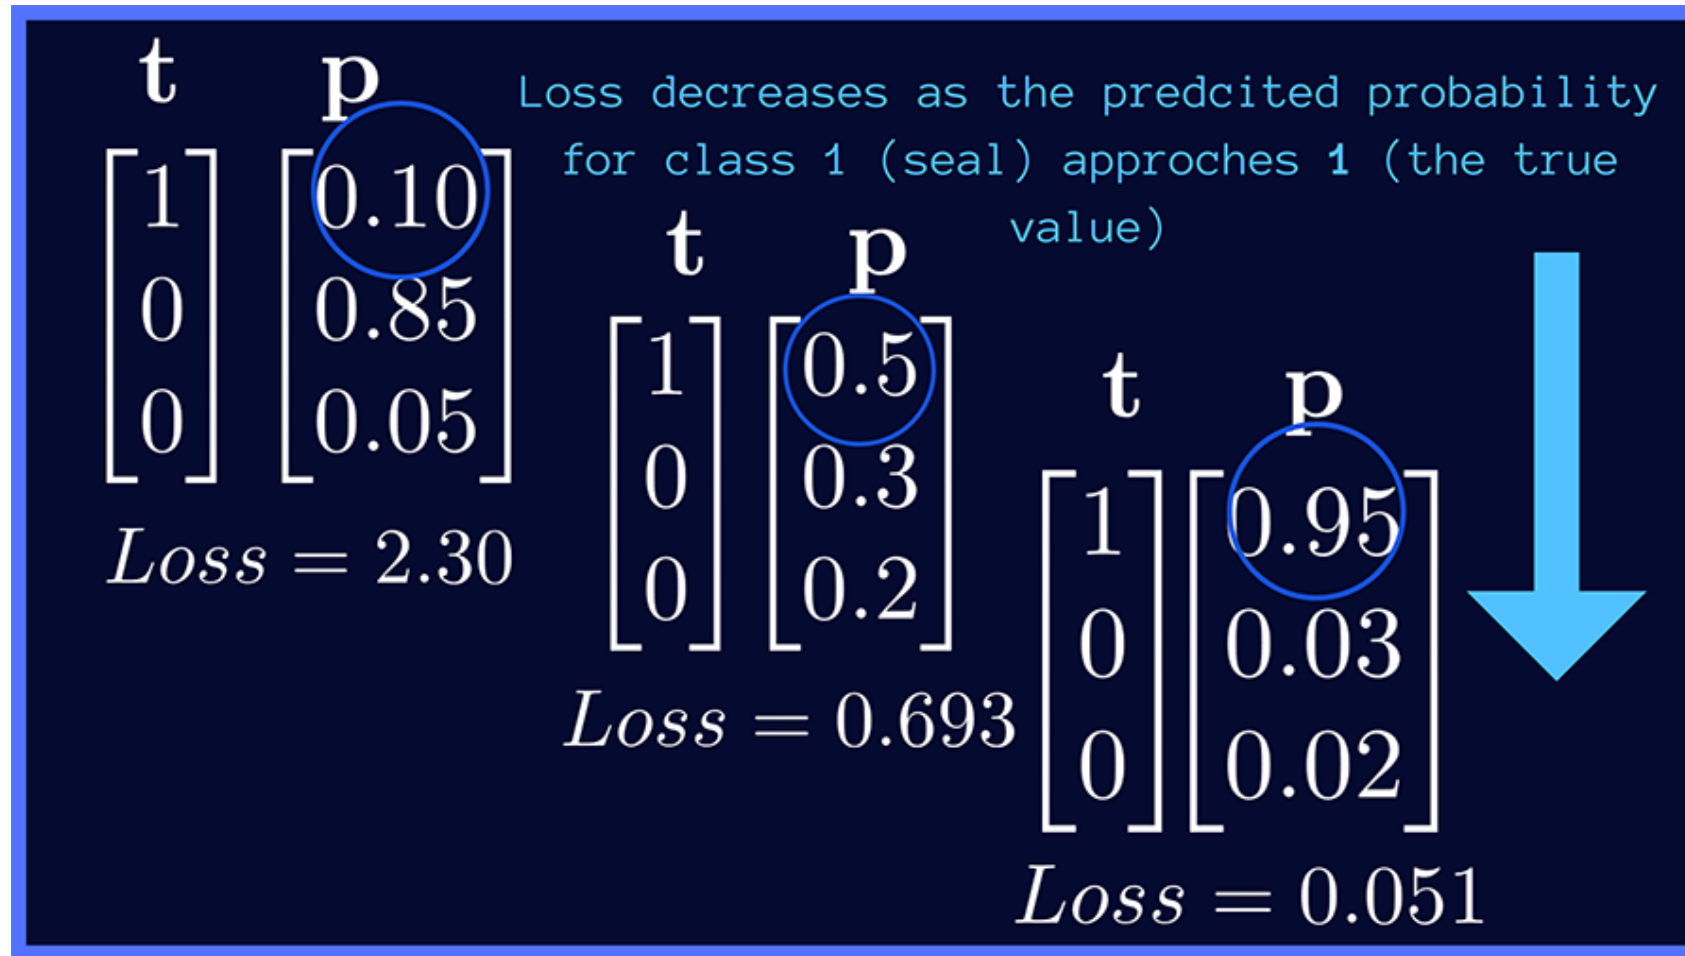
# Multiclass Classification MLP



# Loss function

- We are predicting a *probability distribution*
- Hence we need to use a function which can output
  - a small loss when the probability distributions are very similar
  - a large loss otherwise.
- This is called the ***cross-entropy loss*** and originates from a field called *information theory*.
- We've actually used a version of this loss for logistic regression. It's generally referred to as BinaryCrossEntropy.
- Some Explanations/Resources with the formula:
  - [Entropy, Cross-Entropy and KL Divergence](#) (video resource)
  - [Cross-entropy Loss](#)

# Cross Entropy Loss



# Typical Network Architecture for Classification

Hyperparameter	Binary Classification	Multiclass Classification (also works for 2 classes)
# input neurons	1 per feature	Same
# hidden layers	Depends on the problem. Typically 1-5, but as many as you want is possible.	Same
# neurons per hidden layer	Depends on the problem. Typically 10-100.	Same
# output neurons	1	1 per class
Output Layer Activation	Sigmoid	Softmax
Loss Function	Binary Cross-Entropy	Cross-Entropy