

# Gradient Descent

- Now we know how to evaluate a model, using a cost function, how do we make the model ***learn*** the optimal parameters ?
- In other words, how do we **minimize** the cost function **without testing all** the different possible models ?
- The algorithm used to do this is called ***Gradient Descent***, and is **essential** to most machine learning algorithms, not just linear regression !
- In DL libraries this type of algorithm is called an '**optimizer**' and other variants exist.

# Gradient Descent


- We have some function  $J(\theta_1, \theta_2)$  which we want to minimize...

Outline :

- Start with some initial guess, some random values for  $\theta_1, \theta_2$
- Keep **updating**  $\theta_1, \theta_2$  a little bit to reduce  $J(\theta_1, \theta_2)$  until we end up at a **minimum** (global or local)

# GD intuition

- This is your cost function in 3D
- Imagine you start somewhere near the top of one of the « hills » and your goal is to walk in the direction which will take you down to the bottom the fastest.



# GD formula

$$\begin{array}{l} \text{repeat until convergence } \{ \\ \quad \theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \quad (\text{for } j = 0 \text{ and } j = 1) \\ \} \end{array}$$

- This is the update formula for each of the parameters
- $:=$  signifies assignment
- $\alpha$  is a number called the **learning rate**. If  $\alpha$  is very **large**, then it corresponds to an **aggressive** learning procedure and big steps being taken « downhill » and vice versa.
- $\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$  is a derivative term, which requires a bit of calculus

# GD Intuition


- Why does this update make sense ?
- Why are we putting those 2 terms together ?
- Let's try and get a basic understanding of derivatives before we go any further.

# Derivatives

- The derivative describes **how the output of a function varies** with regard to a very very very **tiny positive nudge** to the **input**, to the point where we consider *almost* no variation in input....
- Informally, the derivative tells you how a function behaves at a particular « instant », i.e. for a given input value.
- The derivative is commonly referred to as « **instantaneous rate of change** »

# Derivatives

- Here is a linear function as an example. What happens when we shift the input by a 'small' value like 0.001
- when  $a=2$  ?
- when  $a=5$  ?



# Derivatives

- With this function, we expect a small positive nudge in the input to make the output increase by 3 times the value of that nudge.

$$f(5.001) = 15.003$$

- In other words the **ratio** between the change in output and the change in input is 3 :

$$\frac{\text{change in } f(a)}{\text{change in } a} = \frac{df(a)}{da} = \frac{0.003}{0.001} = 3$$




# Derivatives


- This is just an example, but formally, the derivative considers this ratio when the input is increased by a **much tinier** amount !
- $\text{Nudge} < 0.000000000000.....1 \Rightarrow$  nudge gets **as close to 0 as possible**
- In this previous example, whatever input value we pick, the derivative will be the same.
- This makes sense since the **function is a line** and the output increases at a constant rate
- Question : What if the derivative was negative everywhere ? What would the function look like ?

# Derivatives

- What if our function isn't a line ?




- The derivative at  $a=2$  is ...



$$\begin{aligned}
 a &= 2 & f(a) &= 4 \\
 a &= 2.001 & f(a) &\approx 4.004 \\
 & & & (4.004004) \\
 \text{slope (derivative) of } f(a) \text{ at} \\
 a &= 2 & \text{ is } & 4. \\
 \frac{d}{da} f(a) &= 4 & \text{ when } a &= 2..
 \end{aligned}$$

- The derivative at  $a=5$  is ...



$$\begin{aligned}
 a &= 2 & f(a) &= 4 \\
 a &= 2.001 & f(a) &\approx 4.004 \\
 & & & (4.004004) \\
 \text{slope (derivative) of } f(a) \text{ at } & & & \\
 a=2 & \text{ is } 4. \\
 \frac{d}{da} f(a) &= \underline{4} \text{ when } a = \underline{2}. \\
 a &= 5 & f(a) &= 25 \\
 a &= 5.001 & f(a) &\approx 25.010 \\
 \frac{d}{da} f(a) &= \underline{10} \text{ when } a = \underline{5}.
 \end{aligned}$$

- Rules exist to compute derivatives
- For example, the function

$$f(a) = a^2$$

$$f'(a) = \frac{d}{da} f(a) = 2a$$

(The notations are called Lagrange and Leibniz notations and are both common)

- If we look at the derivatives/slopes/ratios we calculated previously, this does indeed seem to work !
- Note: the derivative is equal to the **slope of the tangent line** on the graph **at our input value**.

# Derivatives: (optional)

$$f'(x) = \frac{df}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x + \Delta x) - f(x)}{\Delta x}$$

Example 2:  $f(x) = x^2$

$$\begin{aligned} \frac{df}{dx} &= \lim_{\Delta x \rightarrow 0} \frac{f(x + \Delta x) - f(x)}{\Delta x} \\ &= \lim_{\Delta x \rightarrow 0} \frac{x^2 + 2x\Delta x + (\Delta x)^2 - x^2}{\Delta x} \\ &= \lim_{\Delta x \rightarrow 0} \frac{2x\Delta x + (\Delta x)^2}{\Delta x} \\ &= \lim_{\Delta x \rightarrow 0} 2x + \Delta x. \end{aligned}$$


- As  $\Delta x$  approaches 0, the derivative
- Approaches  $2x$ .

# GD Intuition

- Now we have a basic understanding of derivatives, let's apply this understanding to the gradient descent algorithm by using a **simpler example**, with a **cost function of only 1 single parameter**.
- We use  $J(\theta_1)$  instead of  $J(\theta_0, \theta_1)$
- Let's look at a couple scenarios to see how Gradient Descent updates our parameter  $\theta_1$ .

## When the derivative is positive...


- Remember, our cost function looks like a parabola.
- When  $\theta_1$  is too high, we want our optimizer to **reduce** this parameter and bring it closer to the « **sweet spot** », where the cost is minimized.
- Let's see if it does the right thing :





When the derivative is negative...

- When  $\theta_1$  is too low, let's see if Gradient Descent **increases** it and brings it closer to the « sweet spot », where the cost is minimized :



$$\frac{\partial J(\theta_1)}{\partial \theta_1} \leq 0$$
$$\theta_1 := \theta_1 - \alpha (\text{negative number})$$

# Recap

- When the the parameter value is too high, the derivative is positive and the update rule decreases the value for the parameter.


$$\theta_1 := \theta_1 - \alpha \underbrace{\frac{d}{d\theta_1} J(\theta_1)}_{> 0}$$

- Conversely, when the parameter value is too low, the parameter value will be increased by the update rule.

$$\theta_1 := \theta_1 - \alpha \underbrace{\frac{d}{d\theta_1} J(\theta_1)}_{< 0}$$


# Okay so now what about $\alpha$ ?

- Remember the update rule :  $\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1)$
- How does  $\alpha$  influence the update of our parameter  $\theta_1$  ?
- If  $\alpha$  is too small :




# If $\alpha$ is too small

- Many small steps will be taken, which makes Gradient Descent very slow




# If $\alpha$ is too large...

- Gradient descent may « overshoot », go past the minimum. It may even never converge (never find the minimum) and keep jumping around.



If alpha too large




# Question

Suppose  $\theta_1$  is at a local optimum of  $J(\theta_1)$ , such as shown in the figure.

What will one step of gradient descent  $\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1)$  do?

1. Change  $\theta_1$  in a random direction ?
2. Move  $\theta_1$  in the direction of the global minimum of  $J(\theta_1)$  ?
3. Leave  $\theta_1$  unchanged ?
4. Decrease  $\theta_1$  ?



# Recap

- To update our parameter with the Gradient Descent algorithm, we perform 2 essential steps :
    1. Compute the derivative of the parameter with respect to the value we want to minimize (ie. our cost: a score to express how good our model is doing)
    2. Take an optimization step/update the parameter. This update will be proportional to the derivative and the learning rate.
- Large derivative (steep tangent line) + large learning rate = big update



# Piecing everything together

- This is all we need :
  - A **hypothesis** function (our model)
  - A **cost function** (to tell us how well/bad our model is doing)
  - **Gradient Descent or variant** (to update our parameters and get closer to a better model)

Gradient descent algorithm

repeat until convergence {  
     $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$   
    (for  $j = 1$  and  $j = 0$ )  
}

Linear Regression Model

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

# Derivatives vs. Partial derivatives

- Except, instead of having a cost function with a single input, we are back to 2 inputs, our 2 parameters  $\theta_0$  and  $\theta_1$ .
- When we have functions with multiple inputs (known as multivariate functions), computing 1 single derivative is no longer enough!
- The function's « **instantaneous rate of change** » for a given combination of parameters is now determined by 2 values :
  - How does a tiny change in  $\theta_0$  change  $J(\theta_0, \theta_1)$  ?
  - How does a tiny change in  $\theta_1$  change  $J(\theta_0, \theta_1)$  ?

=> Packed together into a vector, these 2 derivatives make up what is referred to as the **gradient**
- Each derivative is a **partial derivative**. (you need both together to get the whole picture !)

# Derivatives vs. Partial derivatives


- *Partial Derivative :*

This comes down to calculating the derivative at each input value, treating the other input as a constant

- We pretend for a second that the other input value has basically no effect on the function
  - when looking at  $\theta_0$  , we treat  $\theta_1$  as a constant
  - when looking at  $\theta_1$  , we treat  $\theta_0$  as a constant

# Partial derivatives visually

- To help illustrate things and relate them to our simple Gradient Descent intuition:



# Gradient Descent

- Each partial derivative tells us how the function behaves (increases/decreases, quickly/slowly, stays constant...) with respect to a single input
- We can then use this information to know if we should increase or decrease each input to get closer to our minimum cost value !
- **Gradient** : the partial derivatives packed together in a vector
- **Descent** : we want to find the cost function's minimum, using the gradient as a source of information to tell us if the cost is increasing/decreasing with respect to each input.

# Update rule

- So we need to figure out the partial derivatives for each parameter !
- the partial derivative of  $J(\theta_0, \theta_1)$  with respect to  $\theta_1$
- the partial derivative of  $J(\theta_0, \theta_1)$  with respect to  $\theta_2$

# Partial derivatives of $J(\theta_1, \theta_2)$

- You can treat these results as being **given**, in order not to go into the details of the derivation.

- General formula

$$\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) = \frac{\partial}{\partial \theta_j} \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i)^2$$

$$\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) = \frac{\partial}{\partial \theta_j} \frac{1}{2m} \sum_{i=1}^m (\theta_0 + \theta_1 x^i - y^i)^2$$

See here for the [MSE derivative](#)

# Partial derivatives of $J(\theta_1, \theta_2)$

- Here are the partial derivatives obtained (take these at face value for now):

$$j = 0 : \quad \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i)$$

$$j = 1 : \quad \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i) x^i$$

- These formulas allow us to compute the partial derivatives for each of the parameters, which we can then plug into our Gradient Descent algorithm.



# Gradient Descent

- We now have formulas to update our parameters !

repeat until convergence {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x^{(i)}$$

}

# Quick recap to put things into perspective

- We have :
- a **model**, which is a line :

$$h(x) = \theta_0 + \theta_1 x$$

- a **cost function**, to tell us how good/bad our model fits the data:


$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$

- **Gradient Descent**, a method to update our parameters so as to minimize the cost function:

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$$


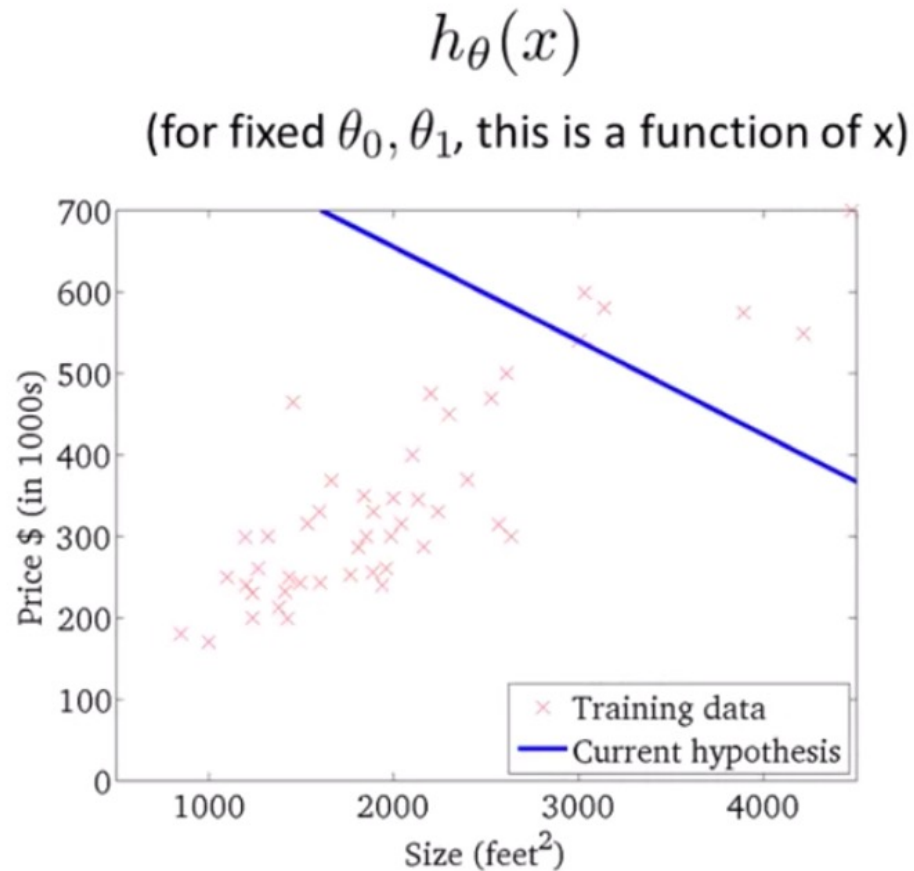
# Update examples

- For linear regression, the cost function will always be bowl-shaped




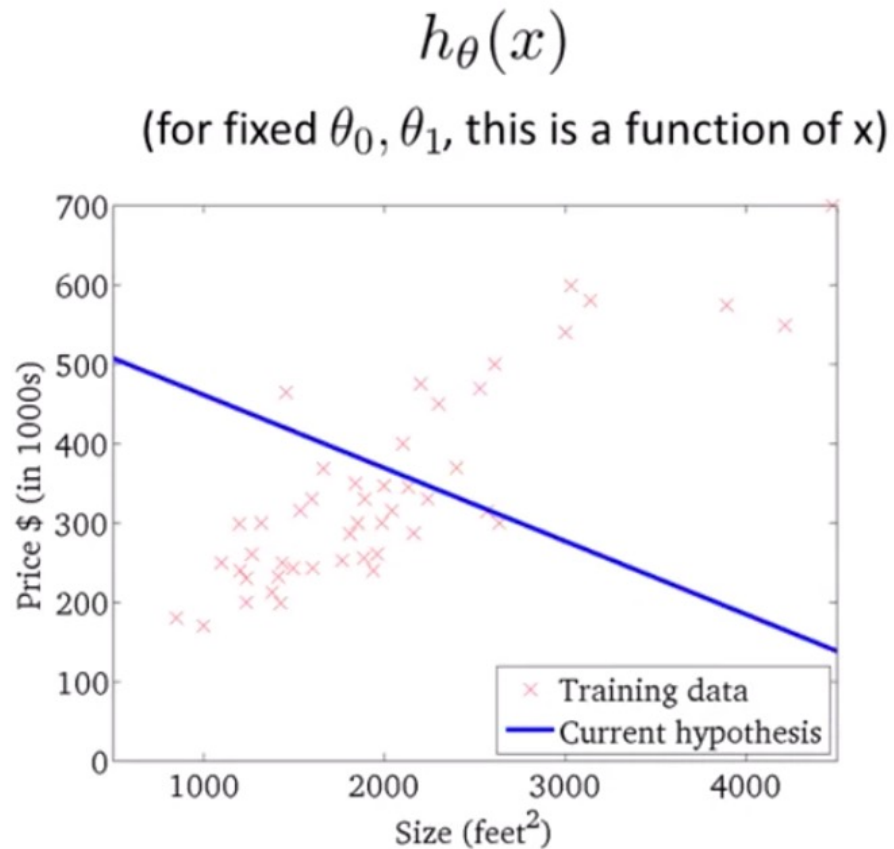
# Update examples

Say we initialize our parameters randomly, this is the model and cost :



# Update examples

As we take Gradient Descent steps, the model (line) seems to be fitting the data better




# Update examples

Until we reach the global minimum

$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )

